



<https://hao-ai-lab.github.io/cse234-w25/>

CSE 234: Data Systems for Machine Learning Winter 2025

LLMSys

Optimizations and Parallelization

MLSys Basics

Today's Learning Goal

- Case study: Matmul on GPU
- Operator Compilation
- **High-level DSL for CUDA: Triton**
- Graph Optimization

Dataflow Graph

Autodiff

Graph Optimization

Parallelization

Runtime

Operator

Triton Programming Model

- Users define **tensors** in **SARM**, and modify them using **torch-like primitives**

Embedded in Python



Kernels are defined in Python using `triton.jit`

Pointer arithmetics



Users construct tensors of pointers and (de)reference them elementwise

Shape Constraints



Must have power-of-two number of elements along each dimension

Example: elementwise add v1 ($z = x + y$)

- Triton kernel will be mapped to a single block (SM) of threads
- Users will be responsible for mapping to multiple blocks

```
import triton.language as tl
Import triton

@triton.jit
def _add(z_ptr, x_ptr, y_ptr, N):
    # same as torch.arange
    offsets = tl.arange(0, 1024)
    # create 1024 pointers to X, Y, Z
    x_ptrs = x_ptr + offsets
    y_ptrs = y_ptr + offsets
    z_ptrs = z_ptr + offsets
    # load 1024 elements of X, Y, Z
    x = tl.load(x_ptrs)
    y = tl.load(y_ptrs)
    # do computations
    z = x + y
    # write-back 1024 elements of X, Y, Z
    tl.store(z_ptrs, z)

N = 1024
x = torch.randn(N, device='cuda')
y = torch.randn(N, device='cuda')
z = torch.randn(N, device='cuda')
grid = (1, )
_add[grid](z, x, y, N)
```

Example: elementwise add v2 ($z = x + y$)

Use multiple blocks

- Index the block and apply offset
- Adds bound check

```
import triton.language as tl
Import triton

@triton.jit
def _add(z_ptr, x_ptr, y_ptr, N):
    # same as torch.arange
    offsets = tl.arange(0, 1024)
    offsets += tl.program_id(0)*1024
    # create 1024 pointers to X, Y, Z
    x_ptrs = x_ptr + offsets
    y_ptrs = y_ptr + offsets
    z_ptrs = z_ptr + offsets
    # load 1024 elements of X, Y, Z
    x = tl.load(x_ptrs, mask=offset<N)
    y = tl.load(y_ptrs, mask=offset<N)
    # do computations
    z = x + y
    # write-back 1024 elements of X, Y, Z
    tl.store(z_ptrs, z)

N = 192311
x = torch.randn(N, device='cuda')
y = torch.randn(N, device='cuda')
z = torch.randn(N, device='cuda')
grid = (triton.cdiv(N, 1024), )
_add[grid](z, x, y, N)
```

Example: elementwise add v2 ($z = x + y$)


- Parametrize block size
- Why we do this?
 - Triton will do tiling for users
 - Avoid manipulating loops

```
import triton.language as tl
Import triton

@triton.jit
def _add(z_ptr, x_ptr, y_ptr, N, BLOCK: tl.constexpr):
    # same as torch.arange
    offsets = tl.arange(0, BLOCK)
    offsets += tl.program_id(0)*BLOCK
    # create 1024 pointers to X, Y, Z
    x_ptrs = x_ptr + offsets
    y_ptrs = y_ptr + offsets
    z_ptrs = z_ptr + offsets
    # load 1024 elements of X, Y, Z
    x = tl.load(x_ptrs, mask=offset<N)
    y = tl.load(y_ptrs, mask=offset<N)
    # do computations
    z = x + y
    # write-back 1024 elements of X, Y, Z
    tl.store(z_ptrs, z)

N = 192311
x = torch.randn(N, device='cuda')
y = torch.randn(N, device='cuda')
z = torch.randn(N, device='cuda')
grid = lambda args: (triton.cdiv(N, args['BLOCK']), )
_add[grid](z, x, y, N)
```

Elementwise Add Performance




Another Example: Softmax

$$y_i = \text{softmax}(\mathbf{x})_i = \frac{e^{x_i}}{\sum e^{x_d}}$$

- How did you implement this in PA1?
 - Think about the potential overhead when compose softmax from primitives
- What if implementing an end-to-end softmax kernel
 - Think about the complexity of implementing in CUDA

Triton Example: softmax



```
import triton.language as tl
Import triton


@triton.jit
def _softmax(z_ptr, x_ptr, stride, N, BLOCK: tl.constexpr):
    # Each program instance normalizes a row
    row = tl.program_id(0)
    cols = tl.arange(0, BLOCK)

    # Load a row of row-major X to SRAM
    x_ptrs = x_ptr + row*stride + cols
    x = tl.load(x_ptrs, mask = cols < N, other = float('-inf'))


    # Normalization in SRAM, in FP32
    x = x.to(tl.float32)
    x = x - tl.max(x, axis=0)
    num = tl.exp(x)
    den = tl.sum(num, axis=0)
    z = num / den;

    # Write-back to HBM
    tl.store(z_ptr + row*stride + cols, z, mask = cols < N)
```

Performance




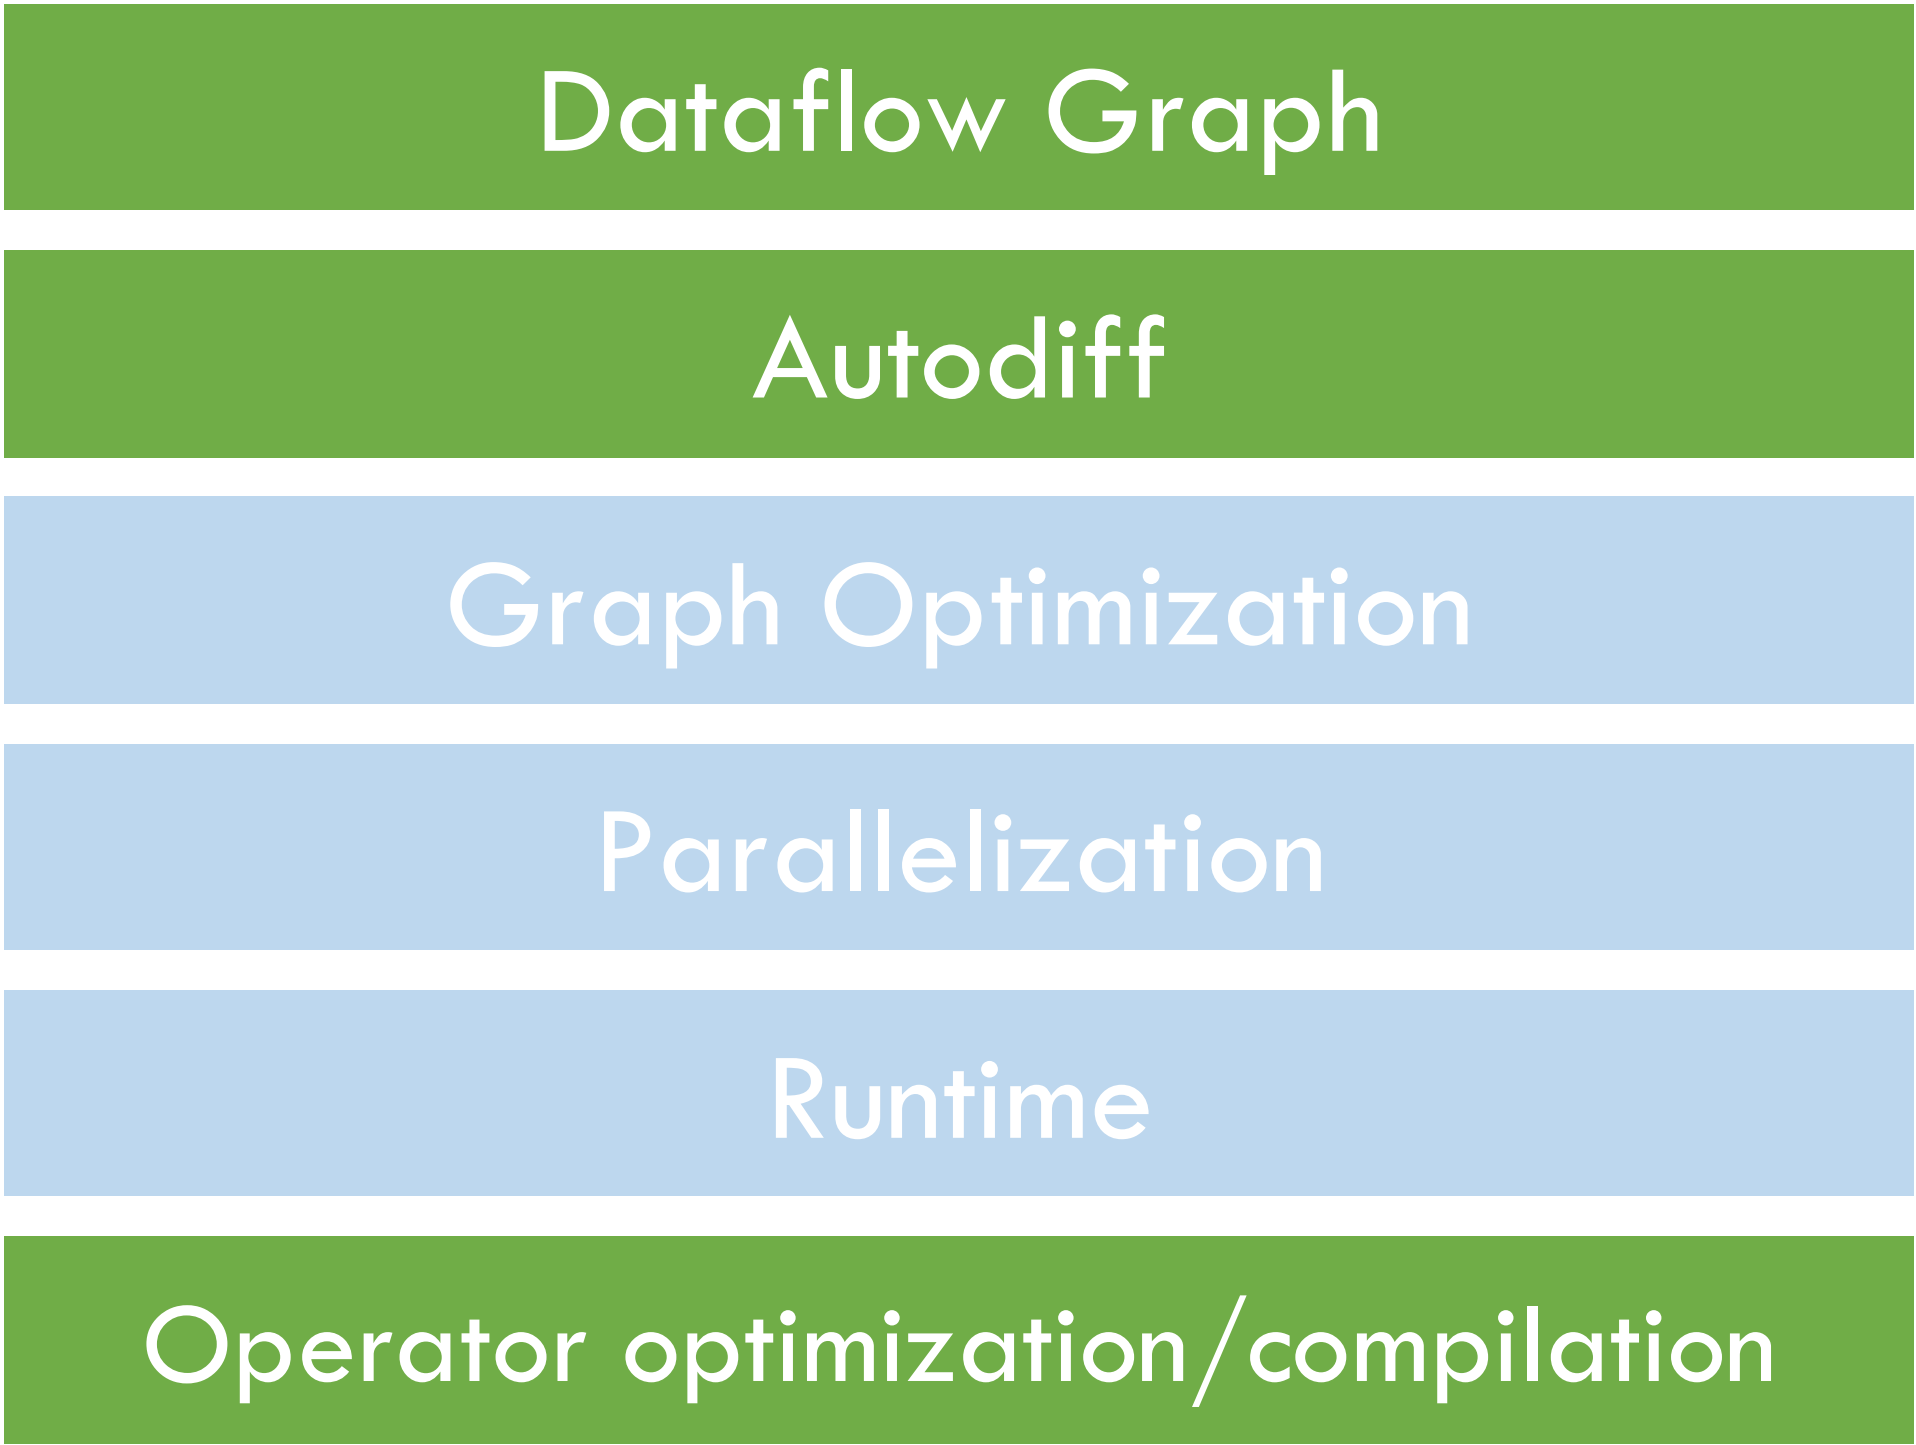
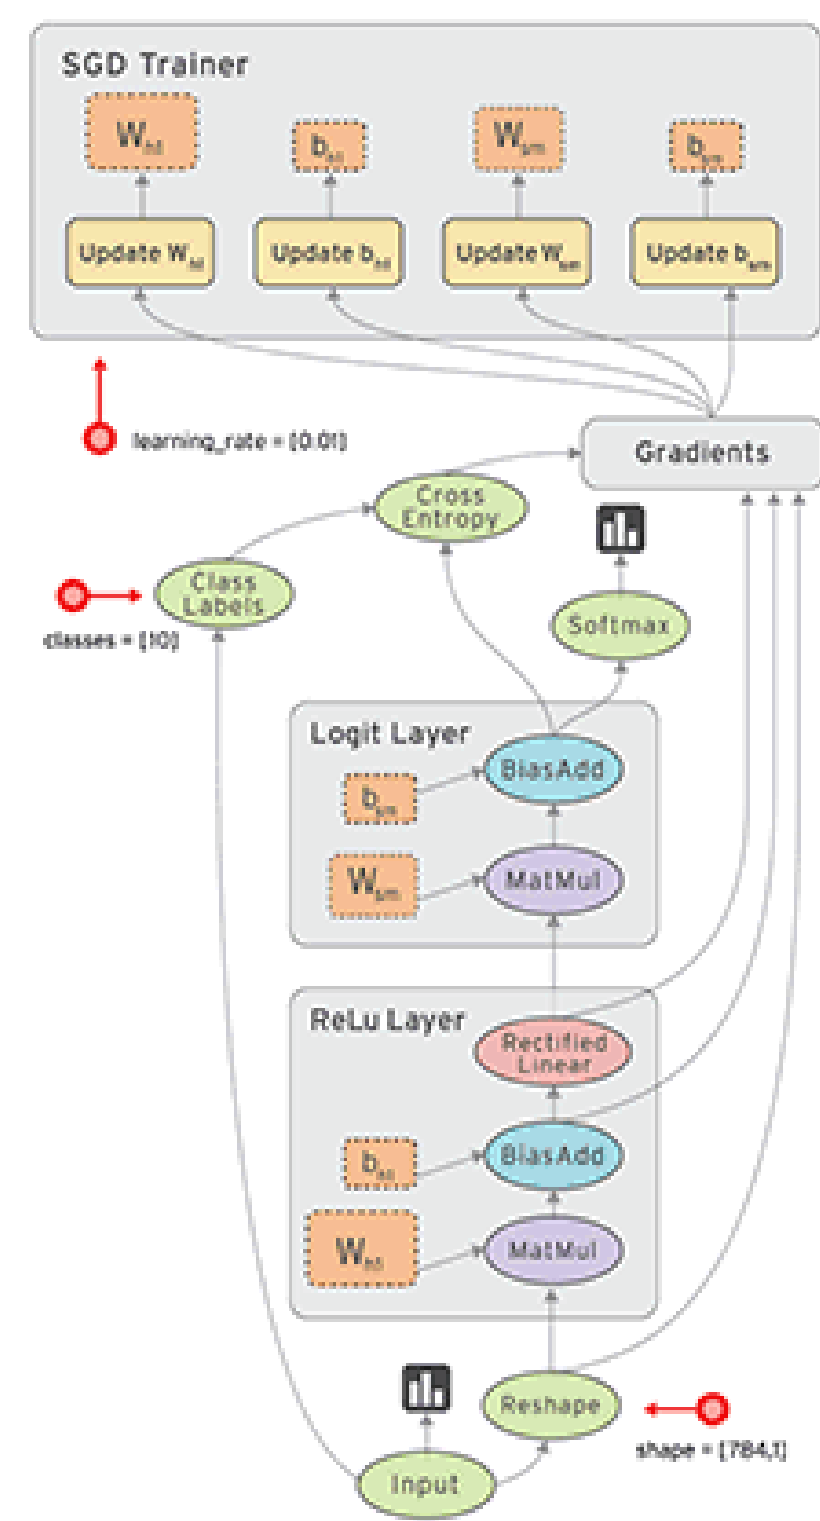
Revisit Triton's Pitch



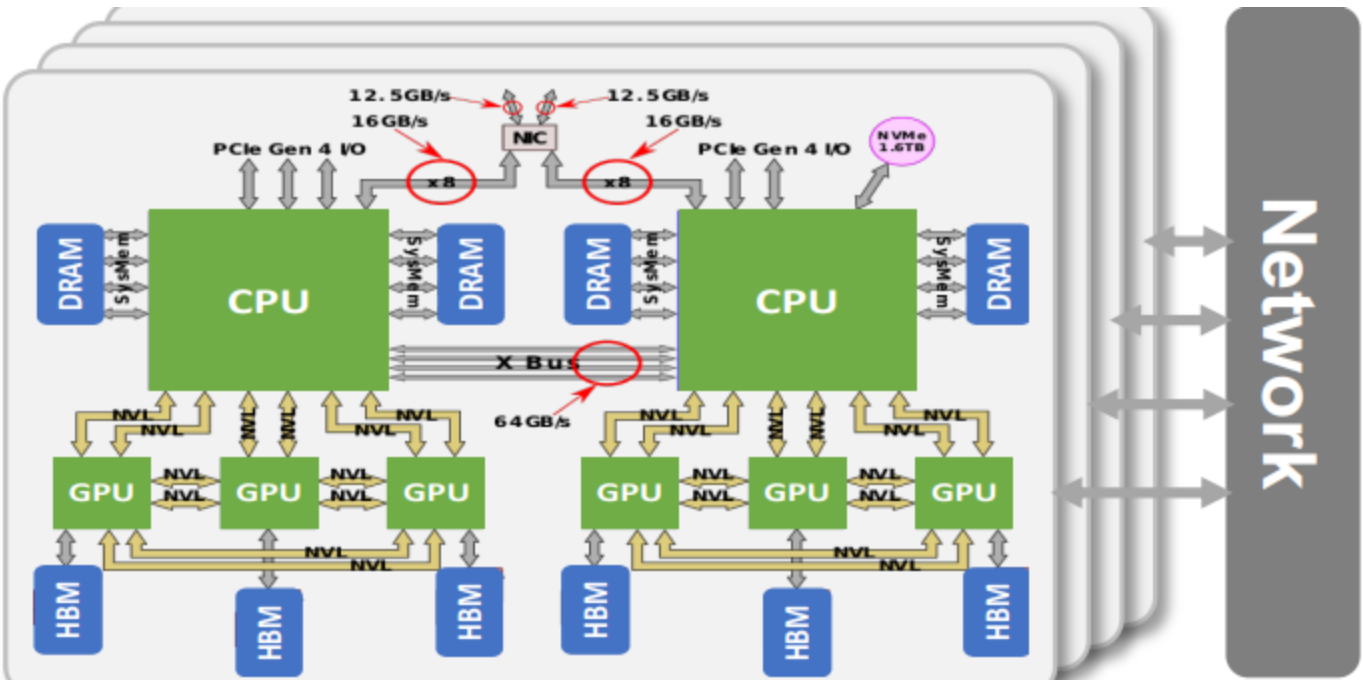
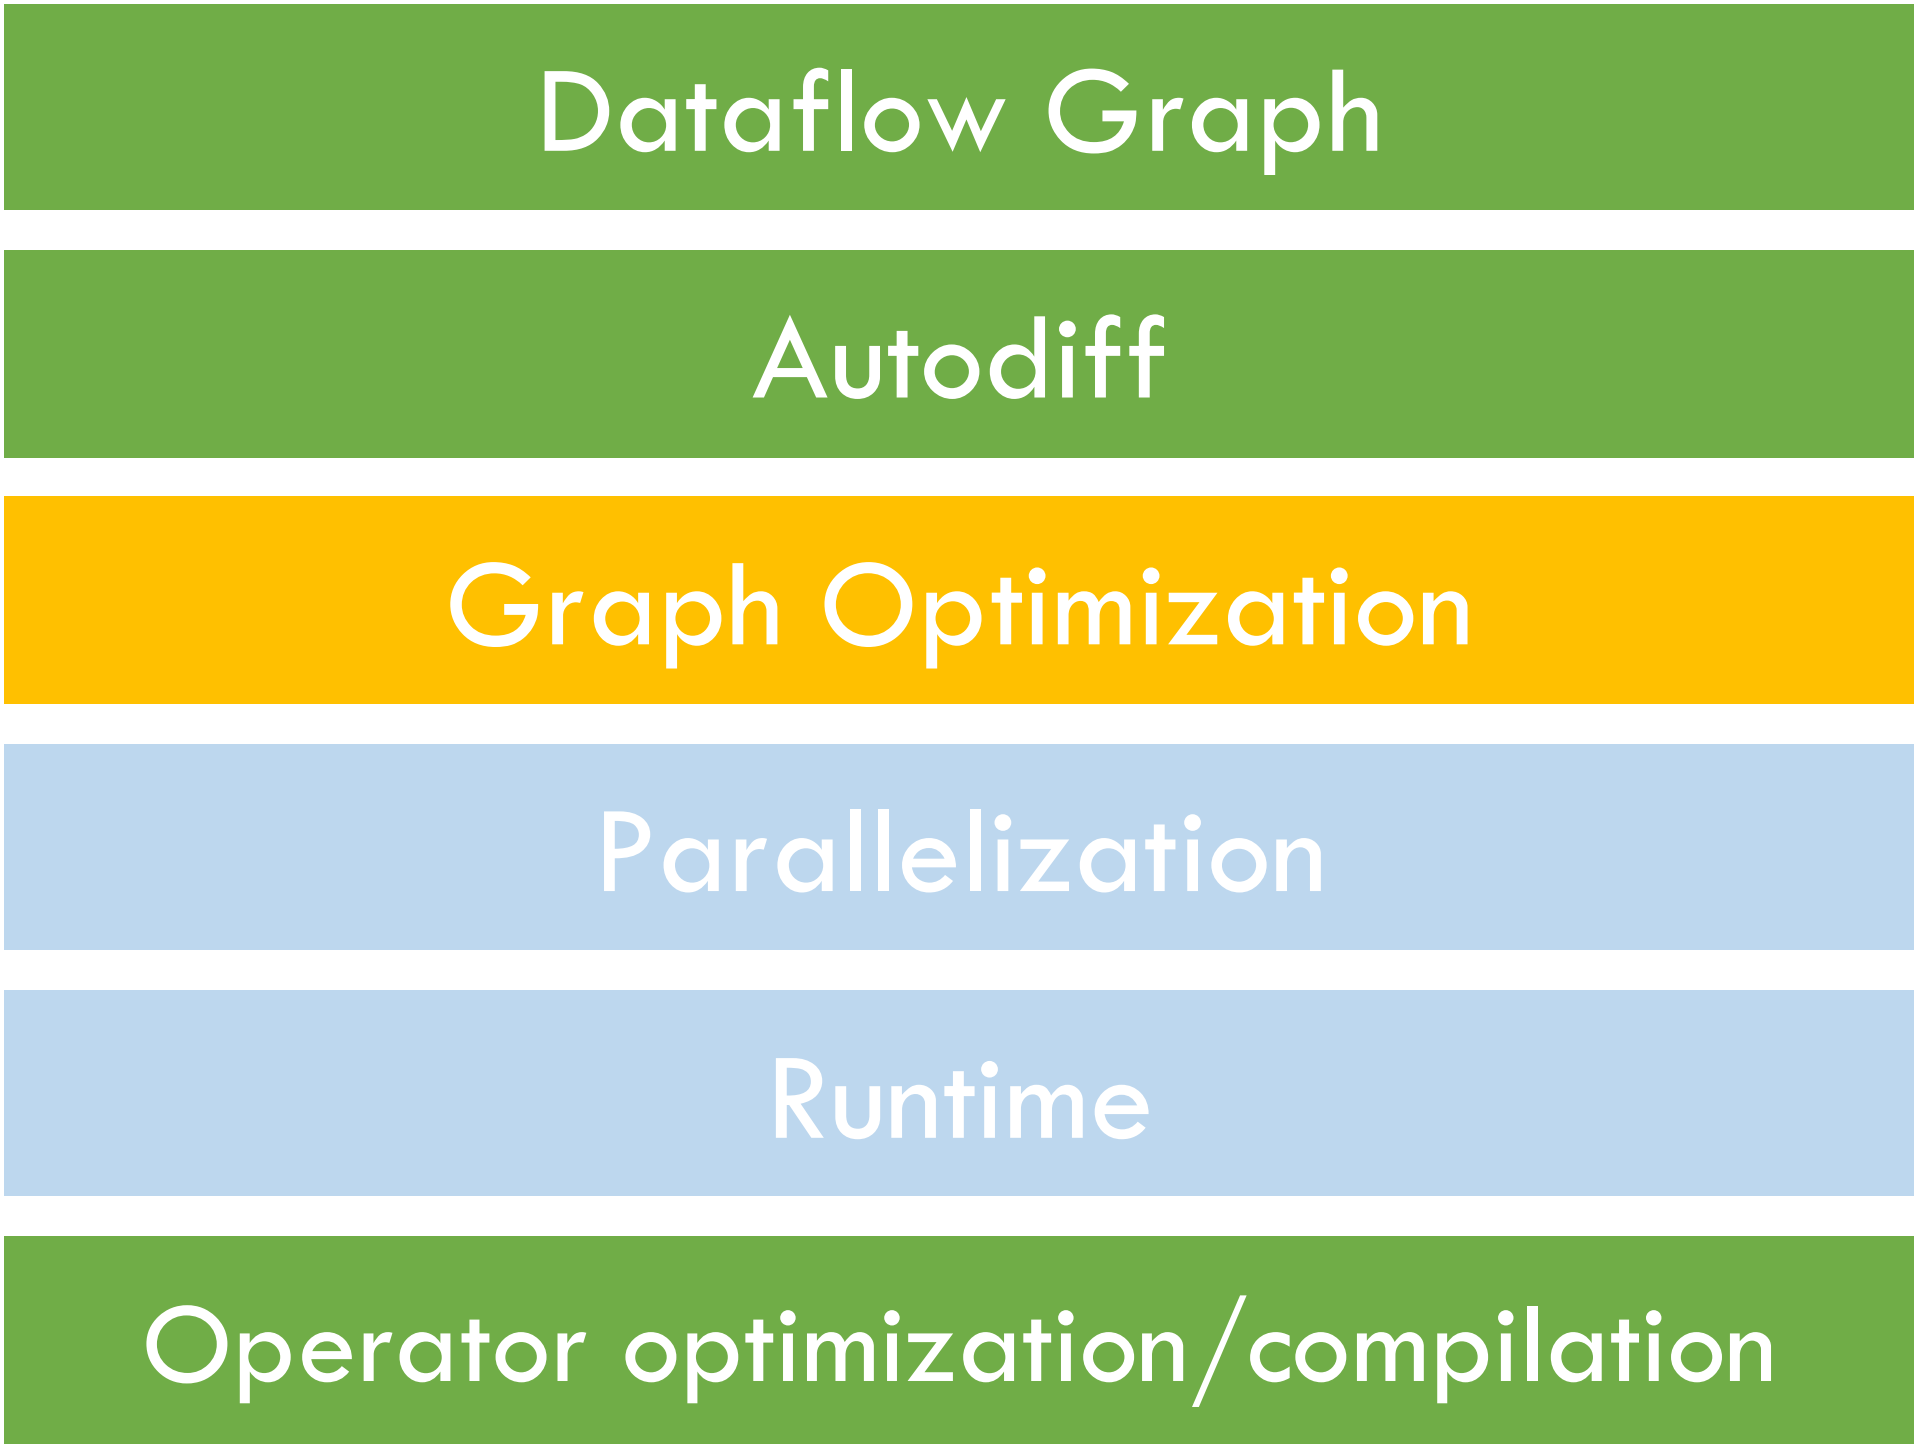
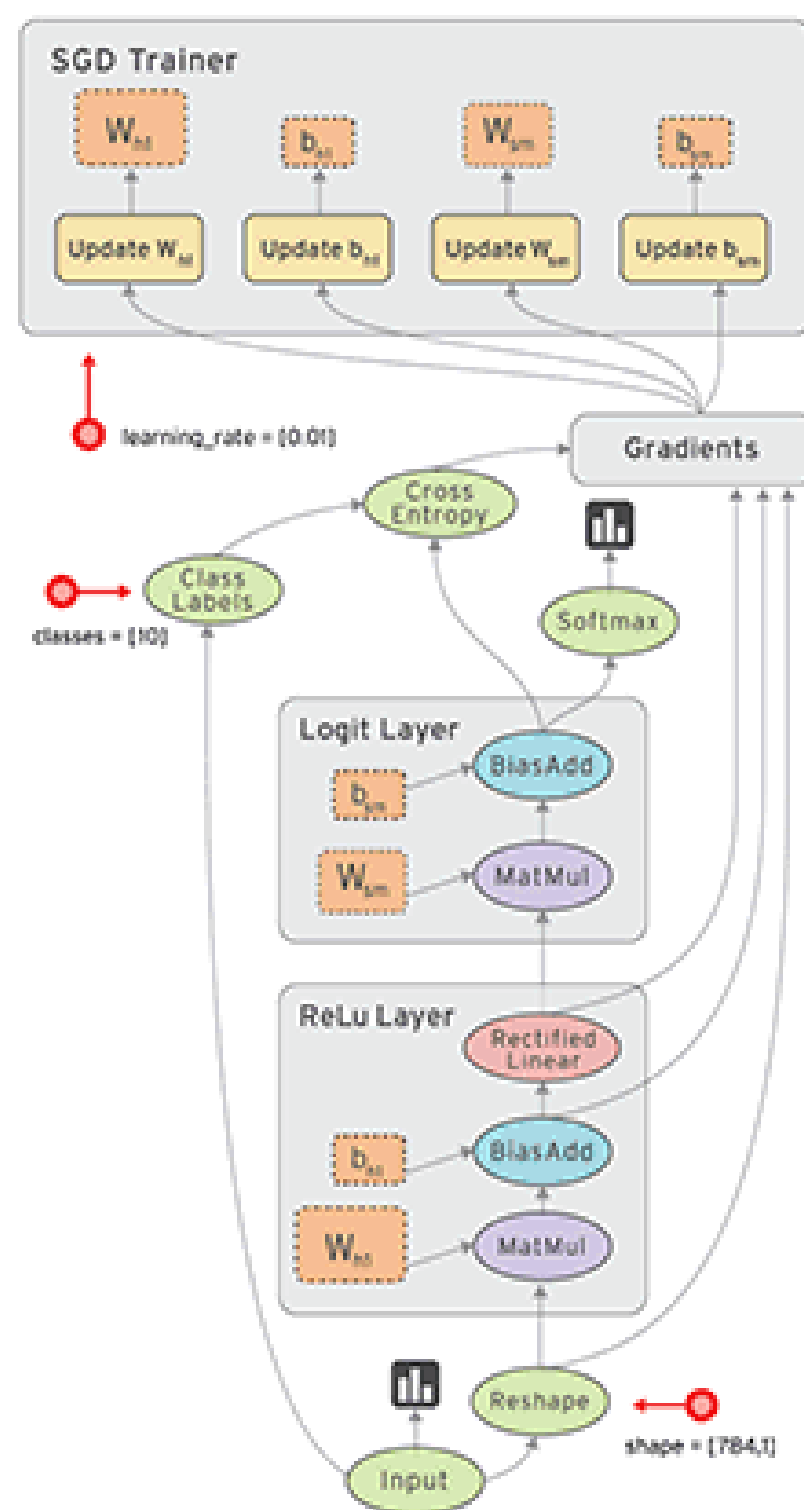
Operator Optimization: Wrapping Up

- Goal: to make individual operator run fast on diverse devices
 1. General ways: vectorization, data layout, etc.
 2. Matmul-specific: tiling to use fast memory
 3. Parallelization SIMD using accelerators
 4. Handcrafted operator kernels vs. automatically compile code
 5. Triton to find the sweet spot

Wrapping Up Operator Optimization





Next: Graph Optimization



Recall Our Goal

- Goal: Rewrite the original Graph G to G' ;
 - G' runs faster than G
 - G' outputs equivalent results
- Straightforward solution: template
 - Human experts write (sub-)graph transformation templates
 - Guarantee correctness and performance gain
 - Run pattern matching over dataflow graph and replace

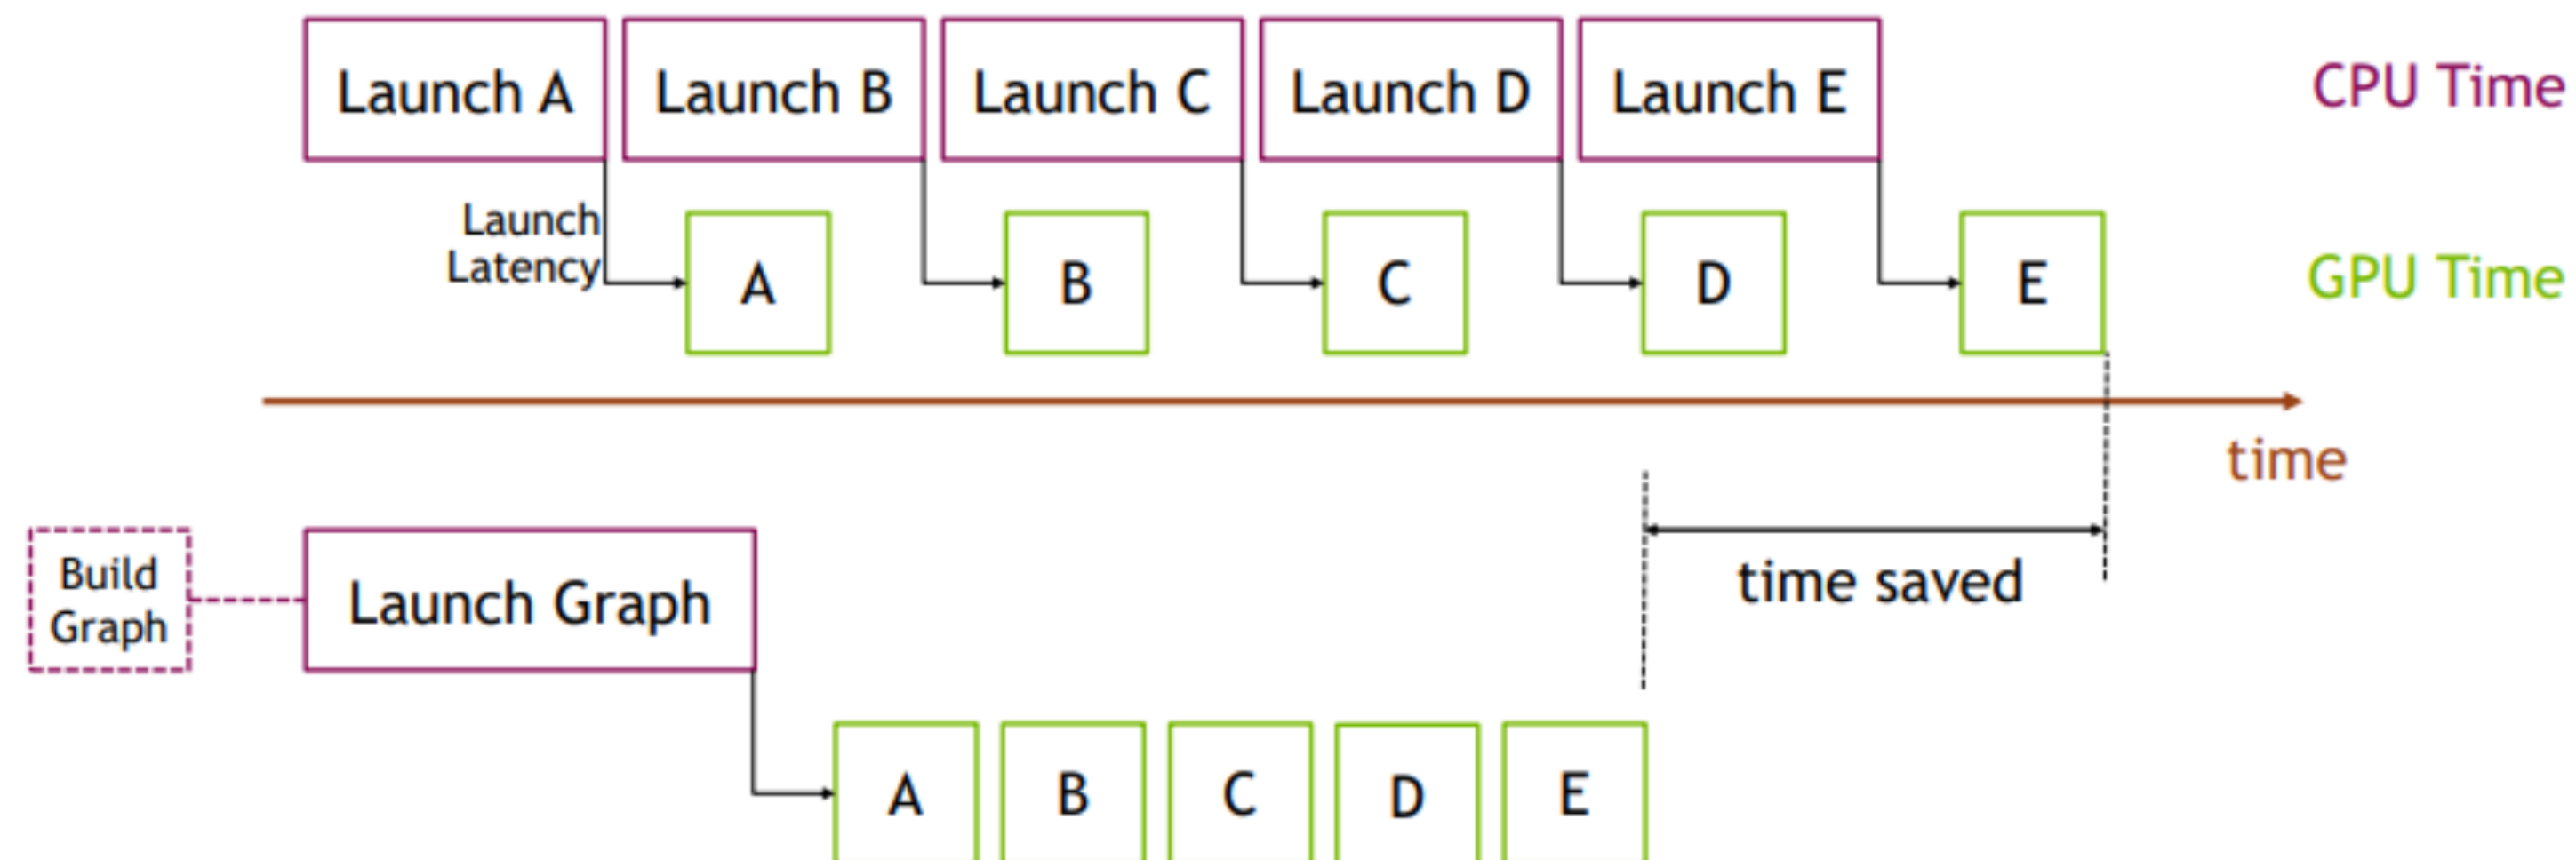
Graph Optimization Templates: Fusion






- Why operator fusion improves performance?
 - Reduce kernel launching
 - Reduce I/O
- Cons:
 - Requiring many fused ops: FusedABCOp
 - At some point, codebase becomes unmanageable

Operator Fusion in Practice: CUDA Graph

- Users are allowed to program using primitives with high-level APIs
- Graph is captured at CUDA level



Graph Optimization Templates: Constant Folding



$A - (-B)$	$A + B$
$A + (A/C1)$	

Common Subexpression Elimination (CSE)

$$\begin{array}{l} \dots \\ c = a + b \\ d = a \\ e = b \\ f = d + e \\ d = x \\ \dots \end{array}$$

$$\begin{array}{l} \dots \\ c^3 = a^1 + b^2 \\ d^1 = a^1 \\ e^2 = b^2 \\ \del{f^3 = d^1 + e^2} \\ f^3 = c^3 \\ d^4 = x^4 \\ \dots \end{array}$$

CSE hit

Dead Code Elimination (DCE)

...
 $c = a + b$
 $d = a$
 $e = b$
 $f = d + e$
 $d = x$
....

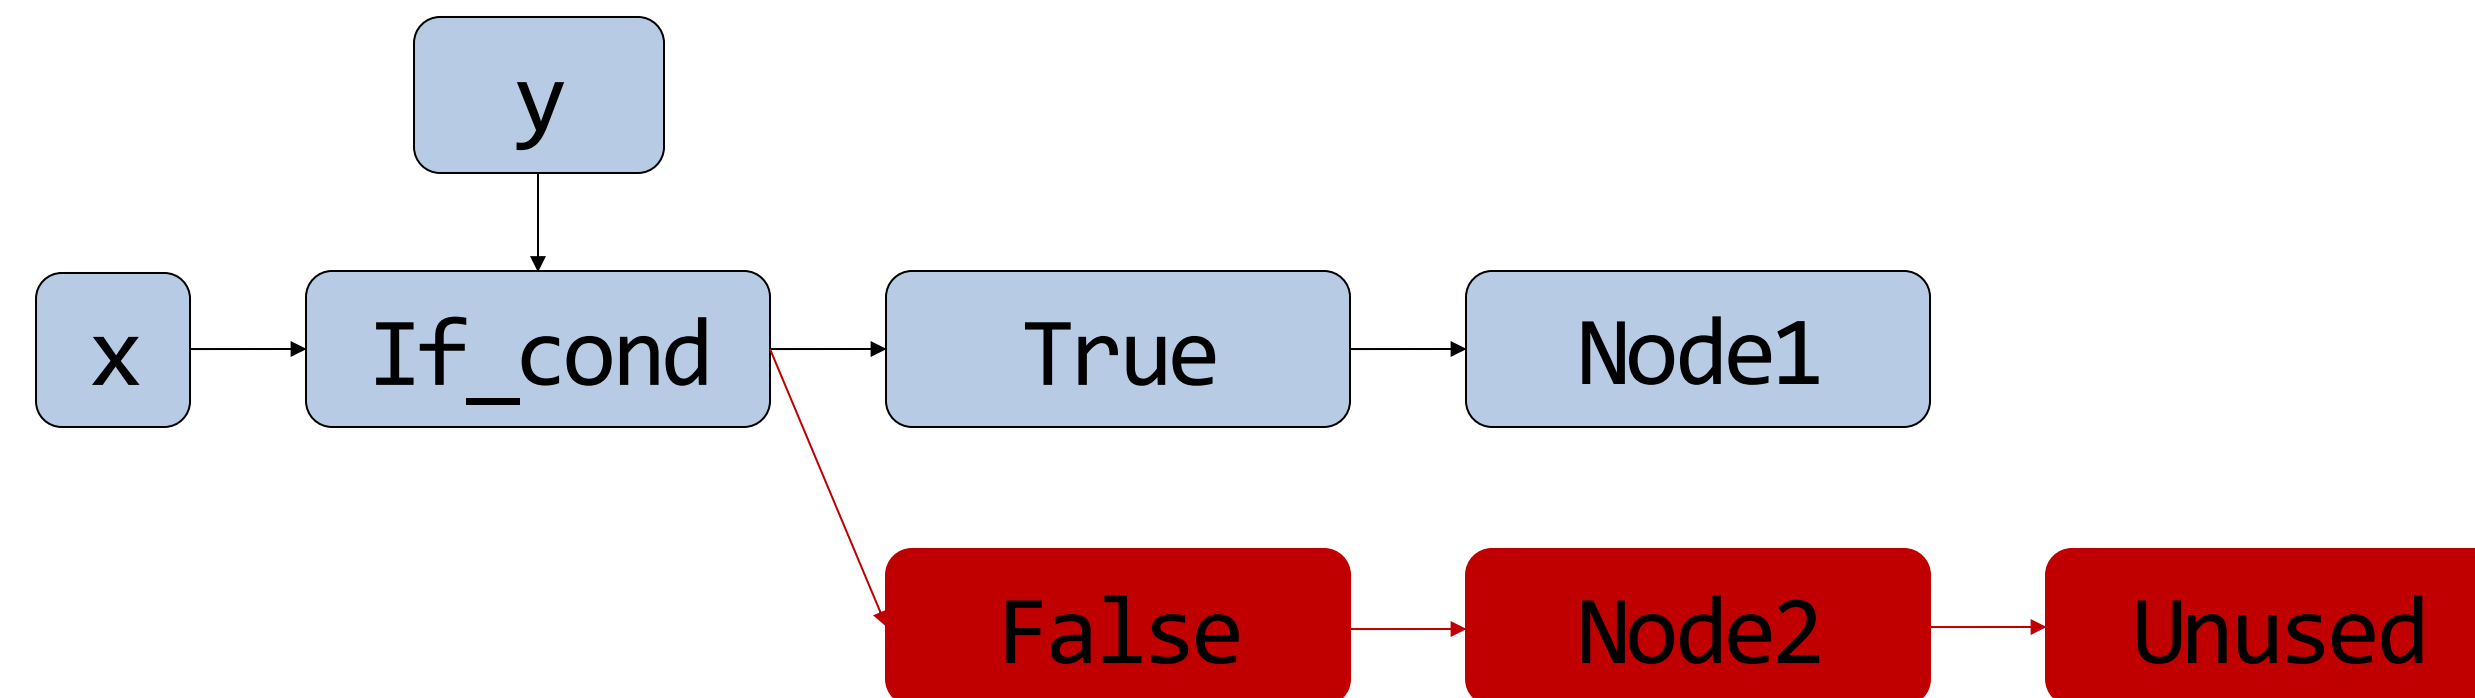

...
 $c^3 = a^1 + b^2$
 $d^1 = a^1$
 $e^2 = b^2$
 ~~$f^3 = d^1 + e^2$~~
 $f^3 = c^3$
 $d^4 = x^4$
....

CSE hit

...
 $c^3 = a^1 + b^2$
 ~~$d^1 = a^1$~~
 $e^2 = b^2$
 ~~$f^3 = d^1 + e^2$~~
 $f^3 = c^3$
 $d^4 = x^4$
....

DCE hit

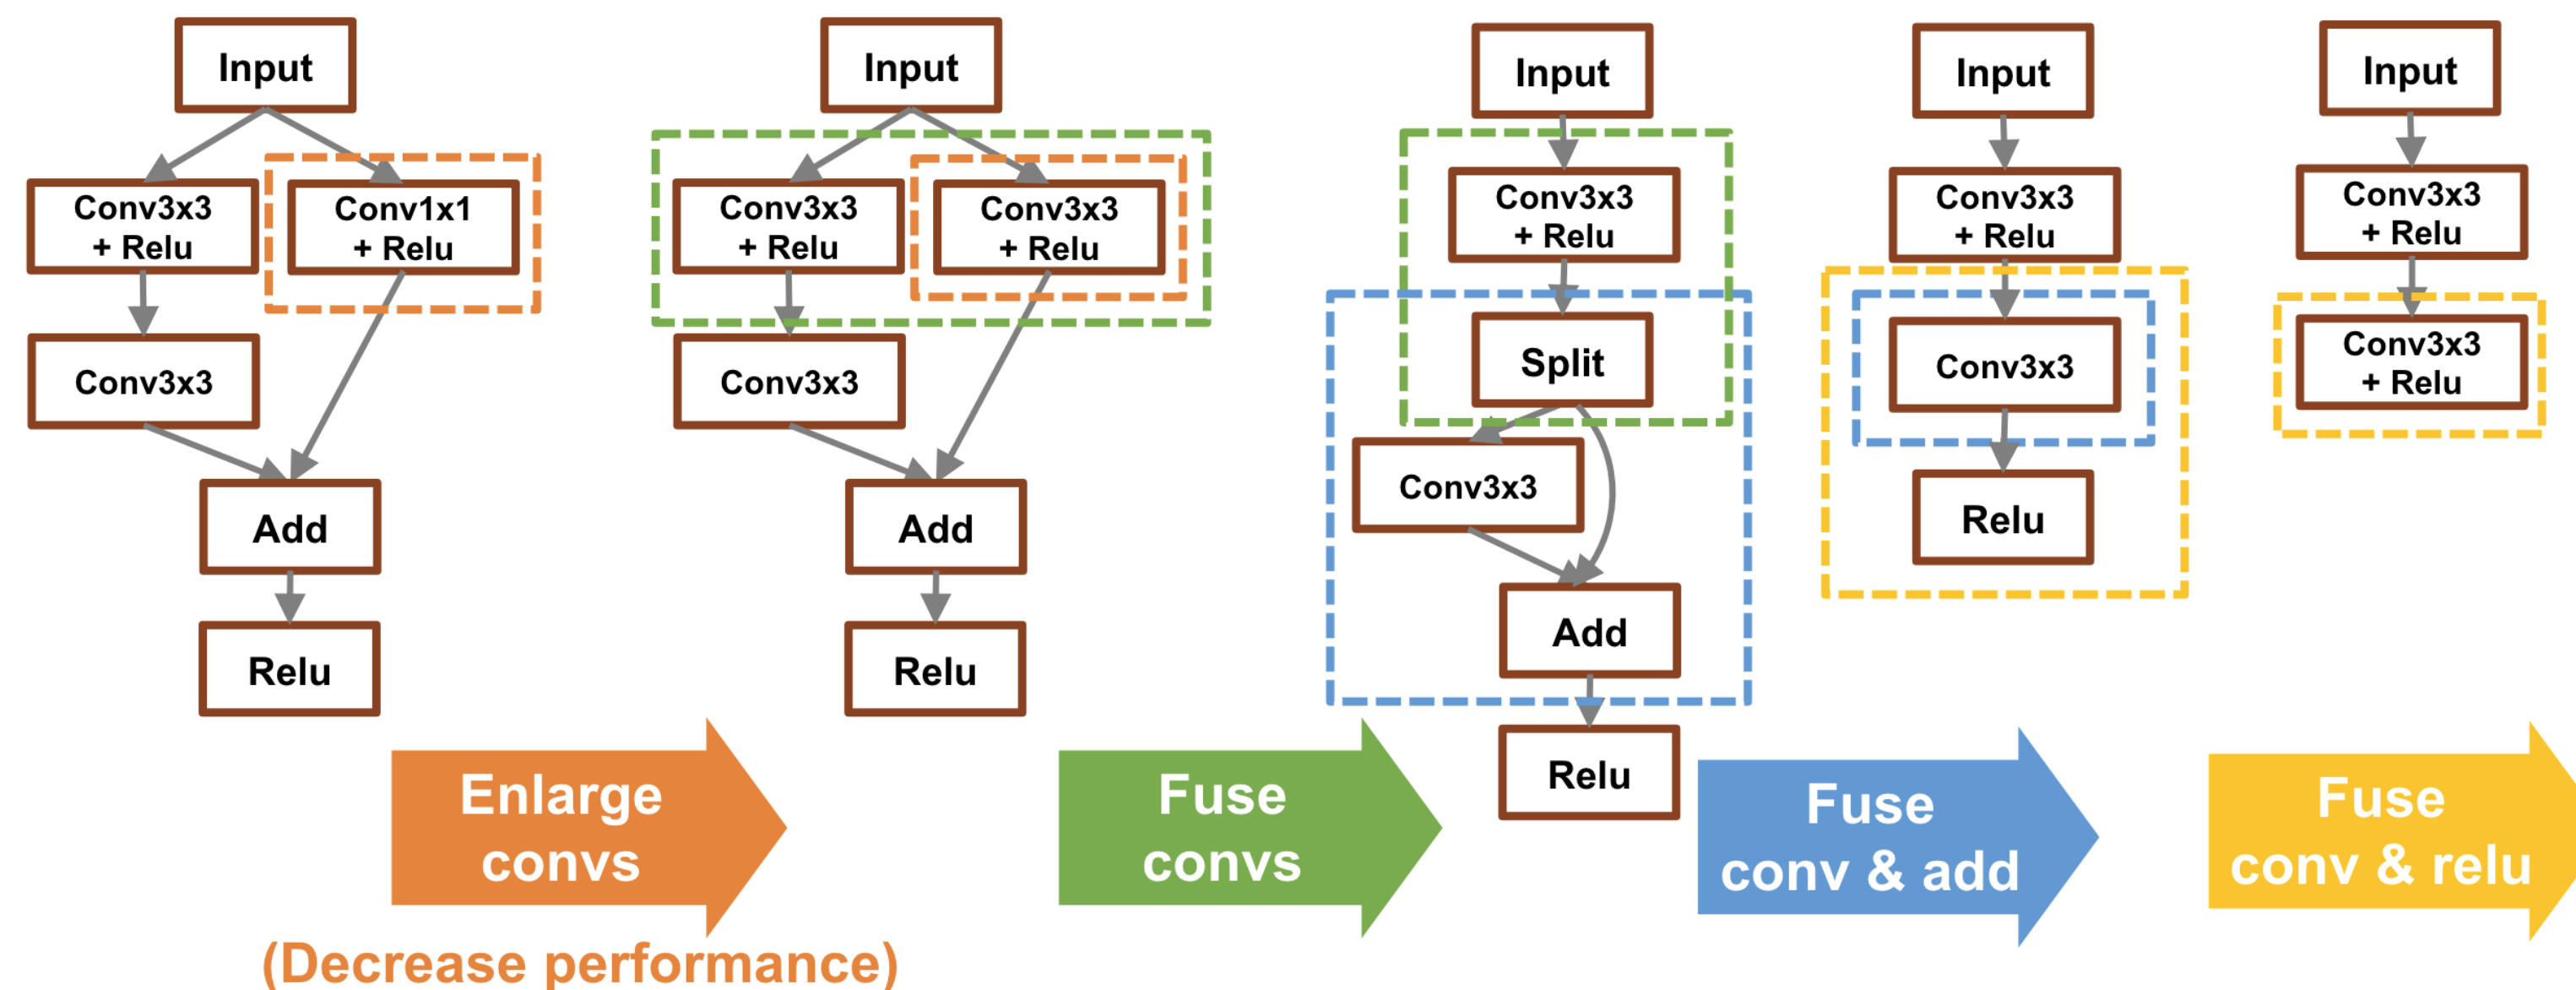
More templates for CSE and DCE



Dataflow Graph
Autodiff
Graph Optimization
Parallelization
Runtime: schedule / memory
Operator

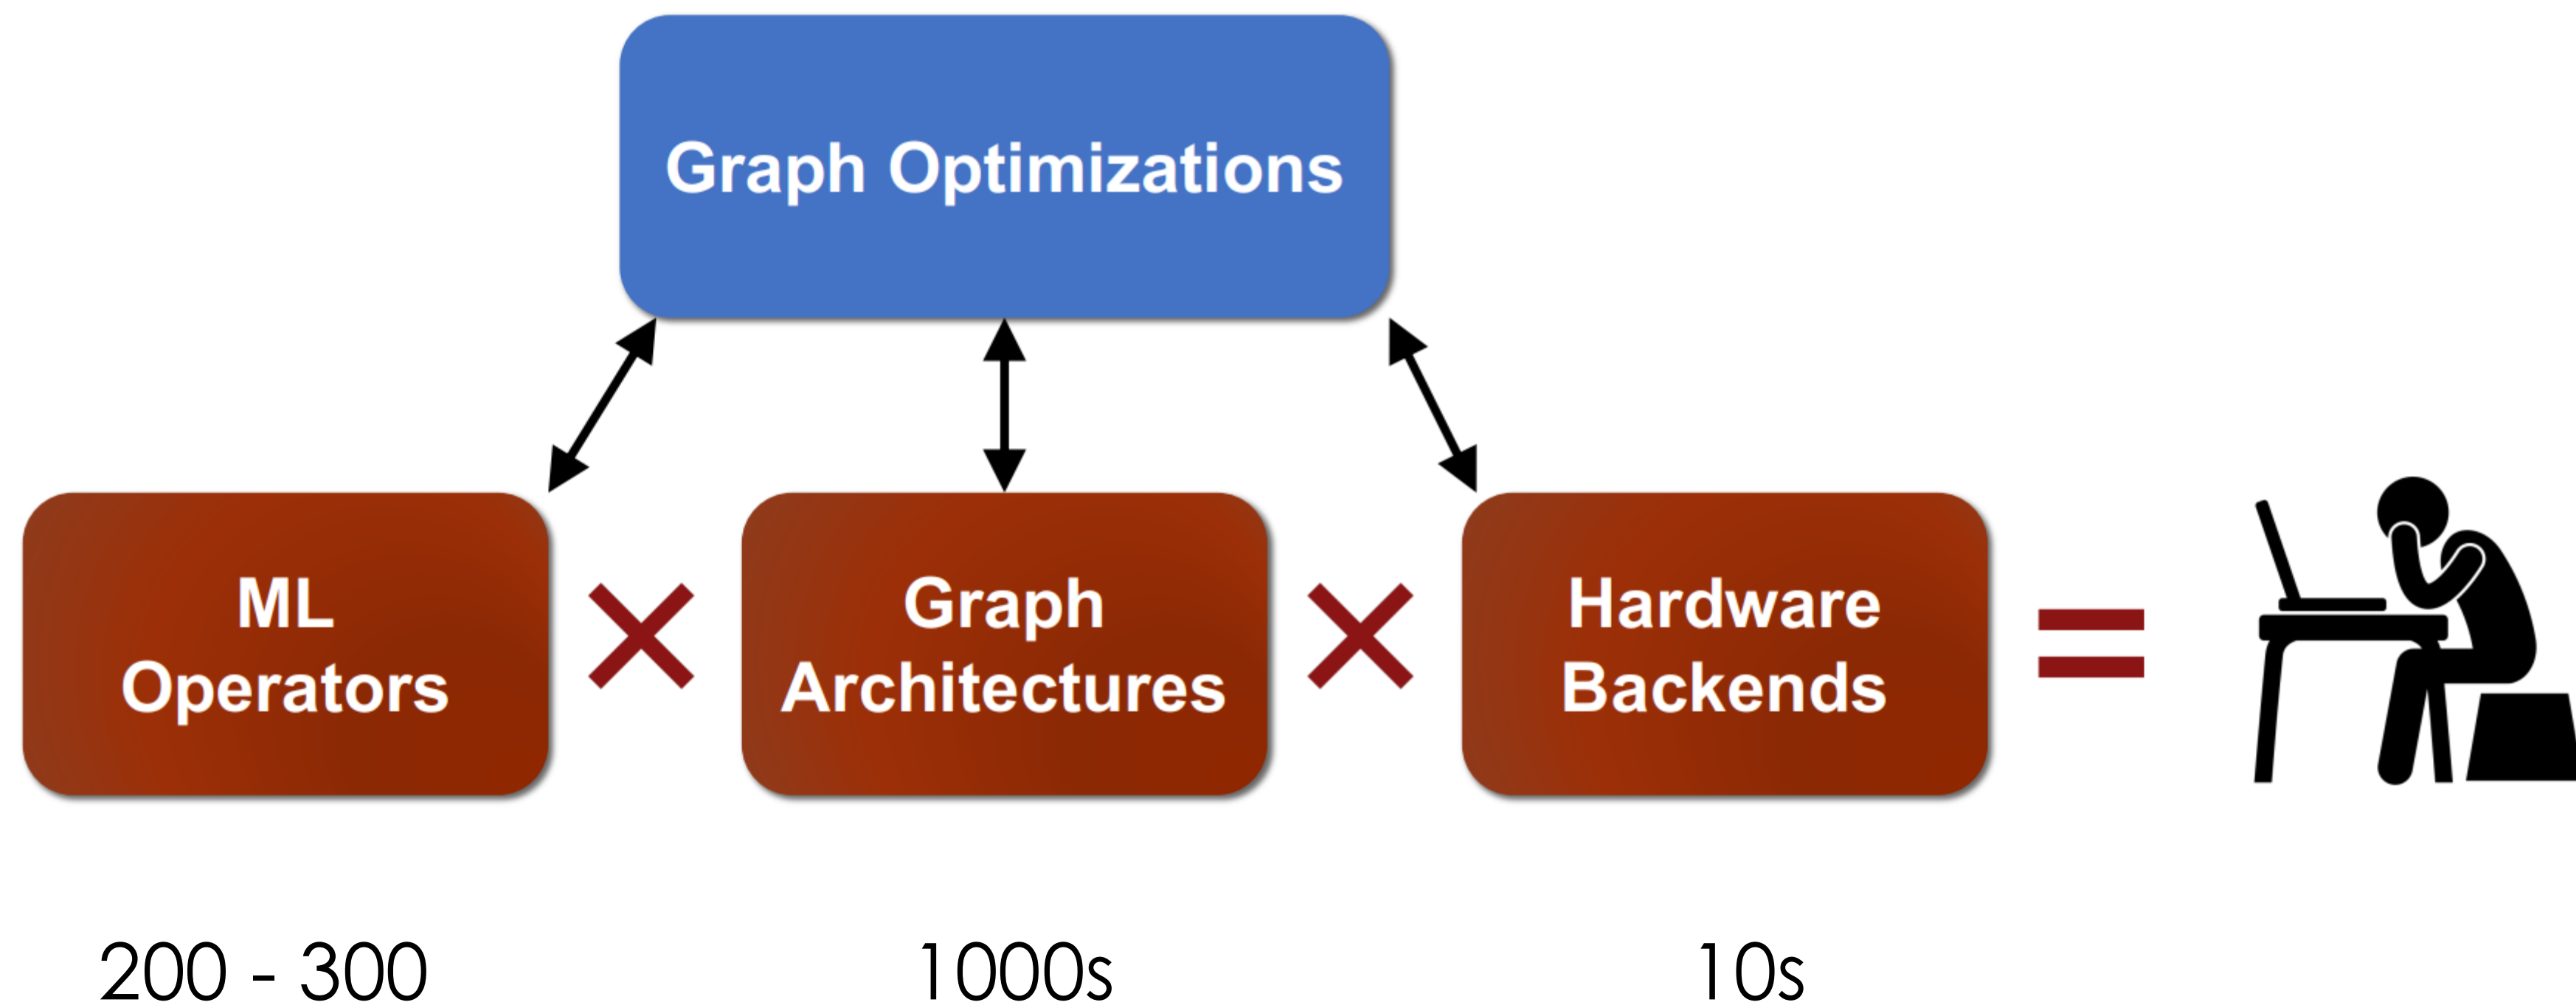
How to ensure performance gain?

- Greedily apply graph optimizations
- Recall the example below



- The final graph is 30% faster on V100 but 10% slower on K80.

Problems of Template-based Graph Optimizations



Problem: Infeasible to manually design graph optimizations for all cases

Problems of Template-based Graph Optimizations

Robustness

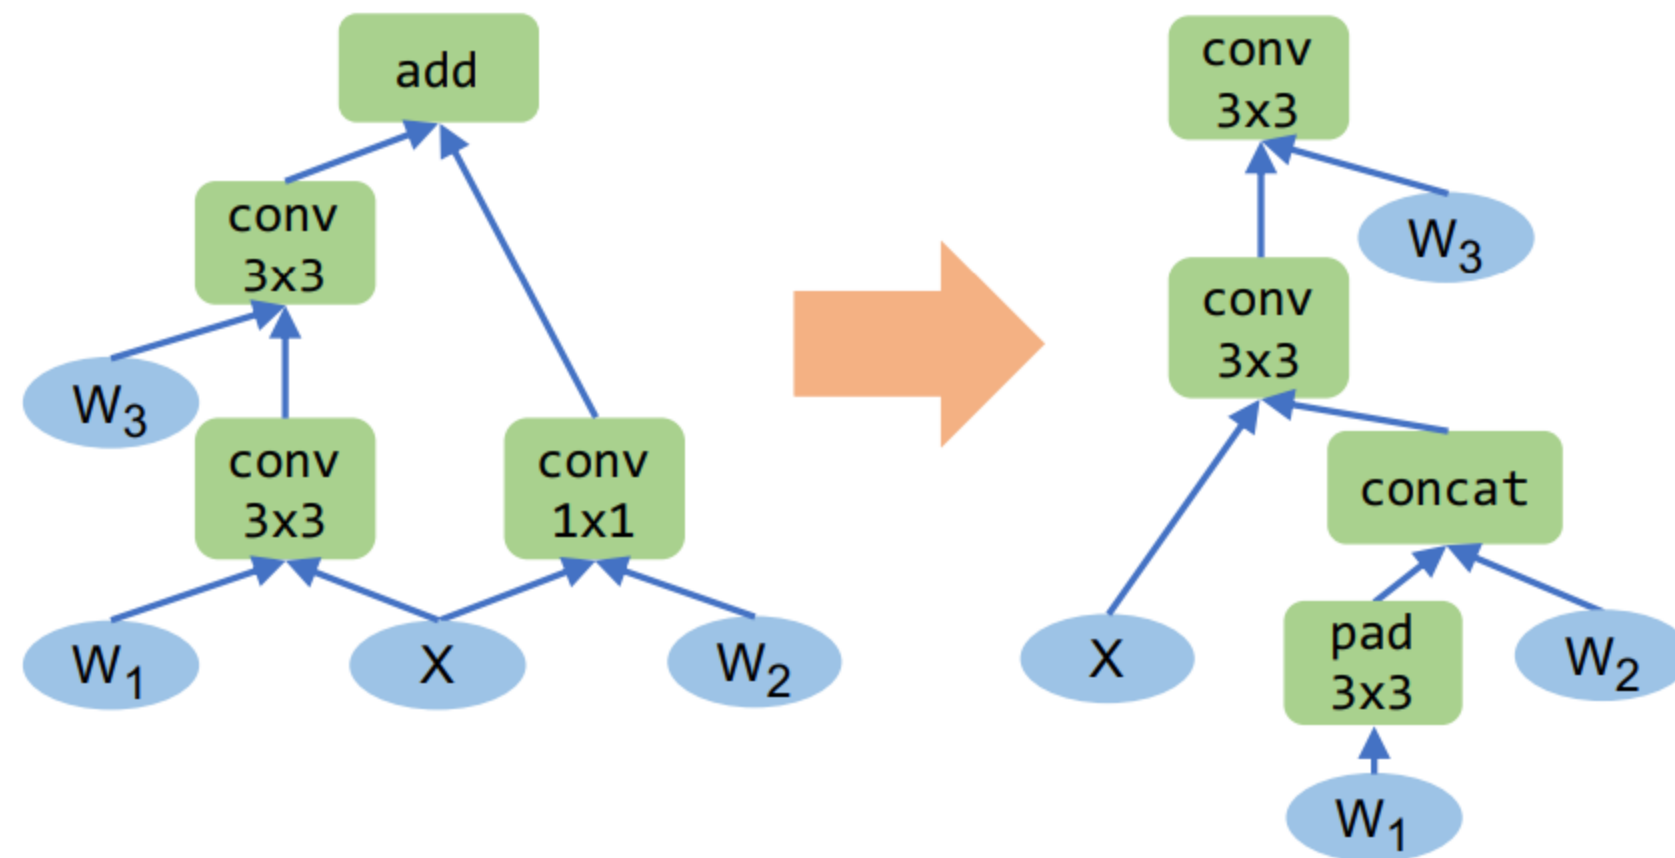
Experts' heuristics do not apply to all DNNs/hardware

Scalability

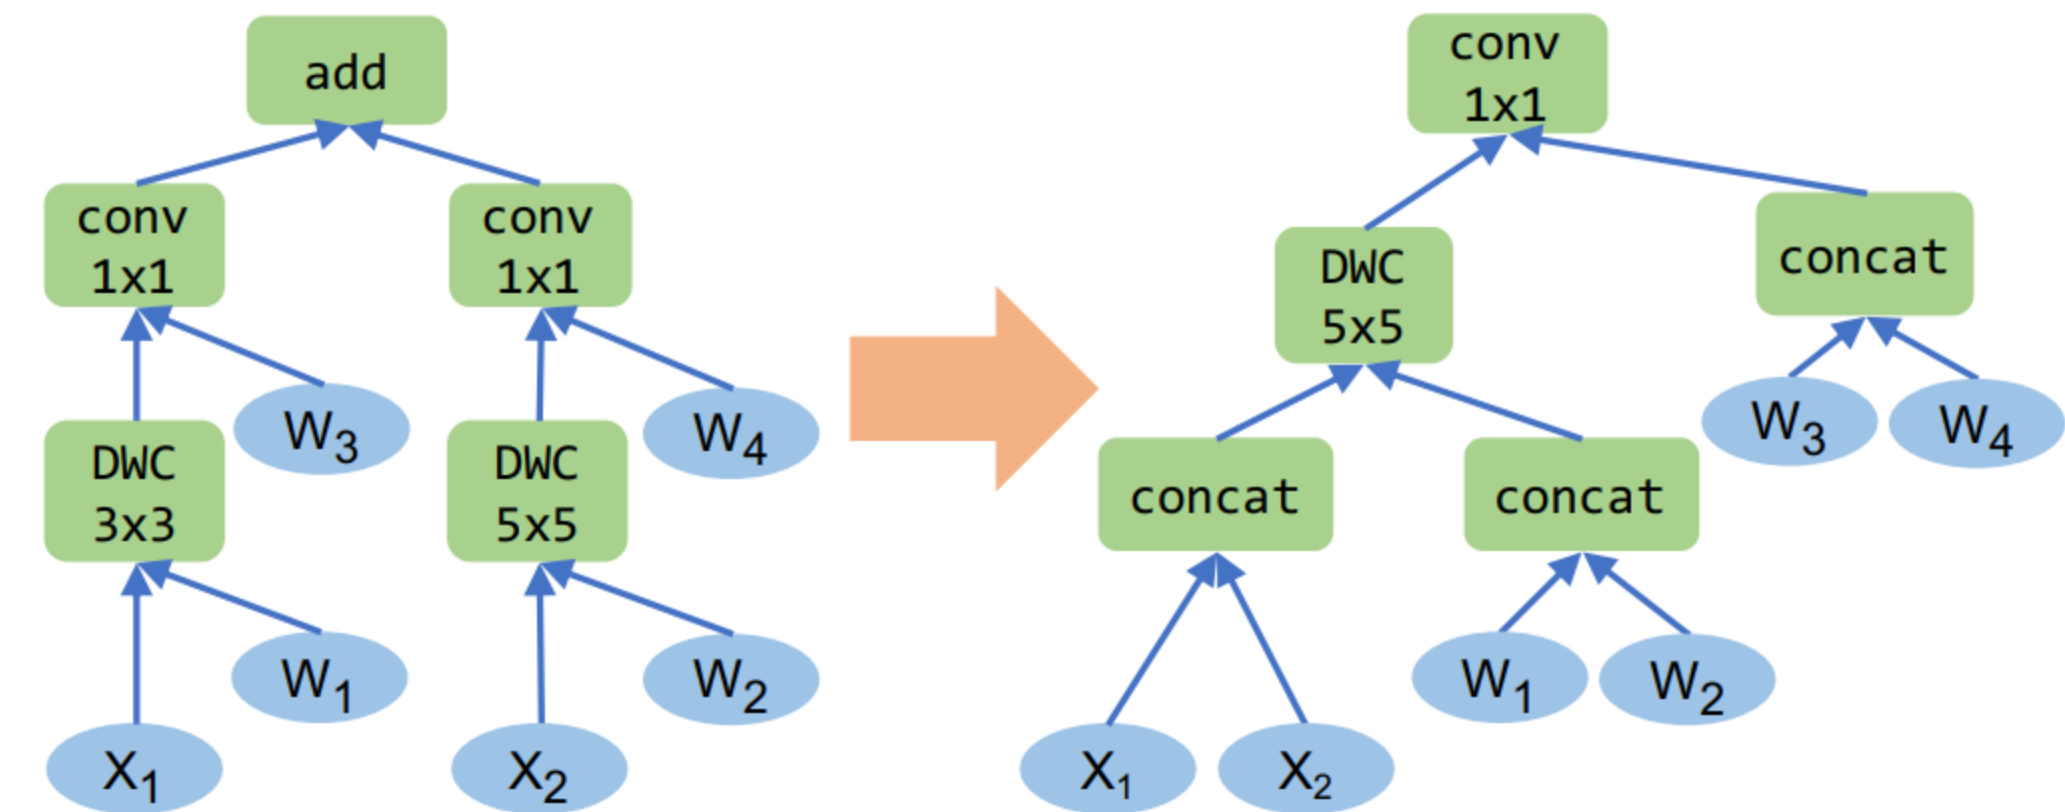
New operators and graph structures require more rules

Performance

Miss subtle optimizations for specific DNNs/hardware



Only apply to **specific hardware**




Only apply to **specialized graph structures**

Automate Graph Transformation

Key idea: replace manually-designed graph optimizations with automated generation and verification of graph substitutions for tensor algebra

Enumerate and Verify ALL possible graph


TASO Workflow



Graph Substitution Generator




Operators supported by
hardware backend



Enumerate all possible graphs up to a
fixed size using available operators


There are many subgraphs even only given 4 Ops

66M graphs with up to 4 operators



A substitution = a pair of equivalent graphs

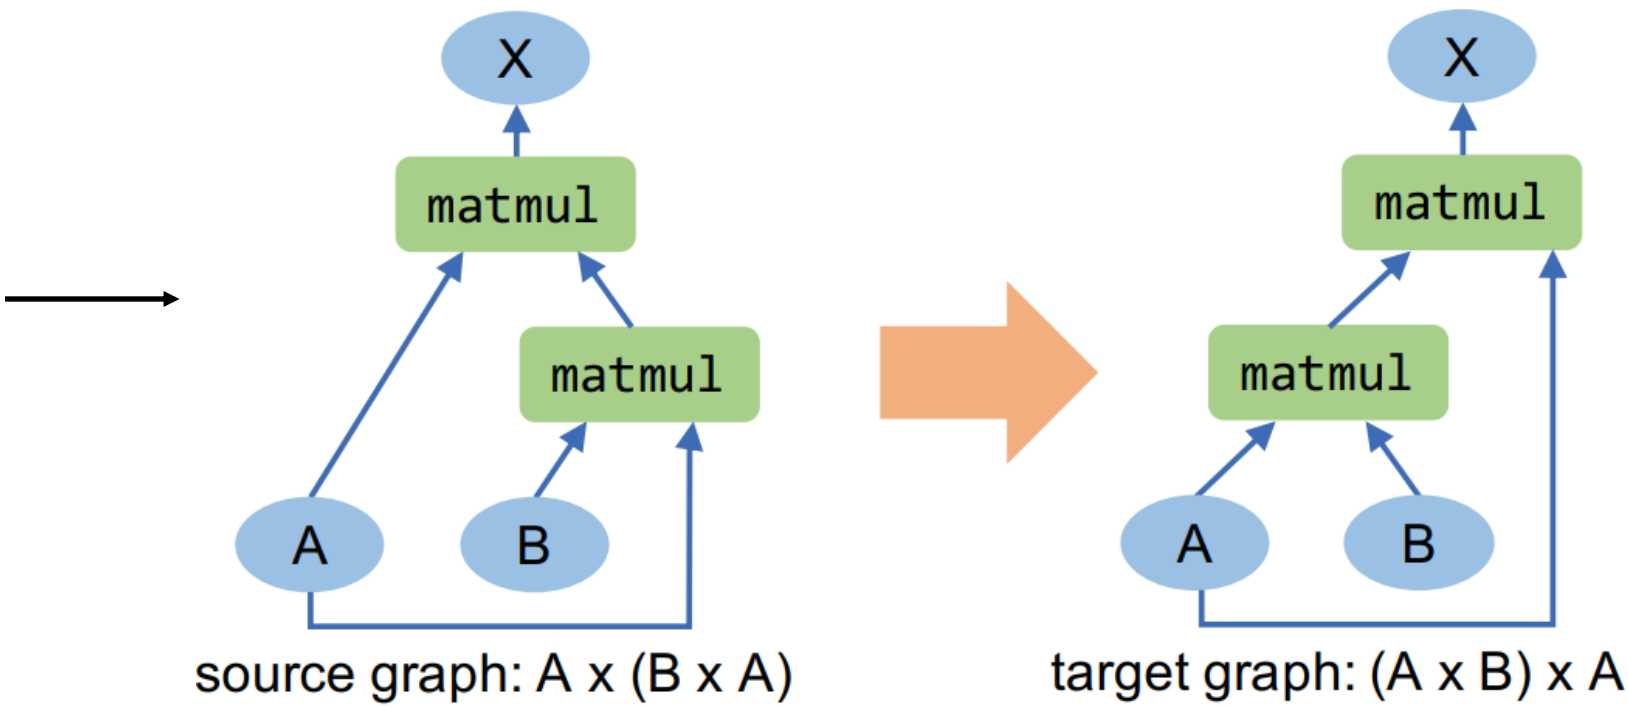
Graph Substitution Generator



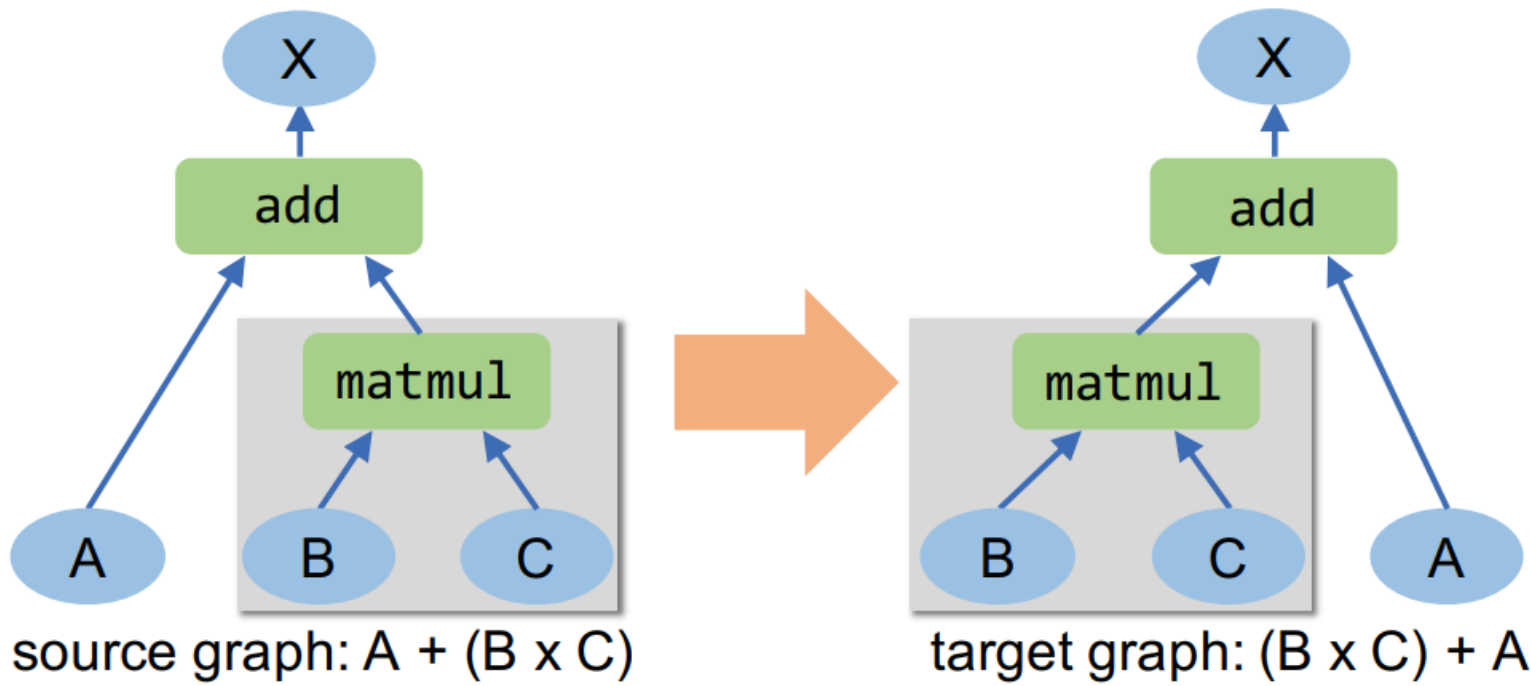
We can generate 28744 substitutions by enumerating graphs with up to 4 ops

Pruning repeated graphs

28744
substitutions



Variable renaming



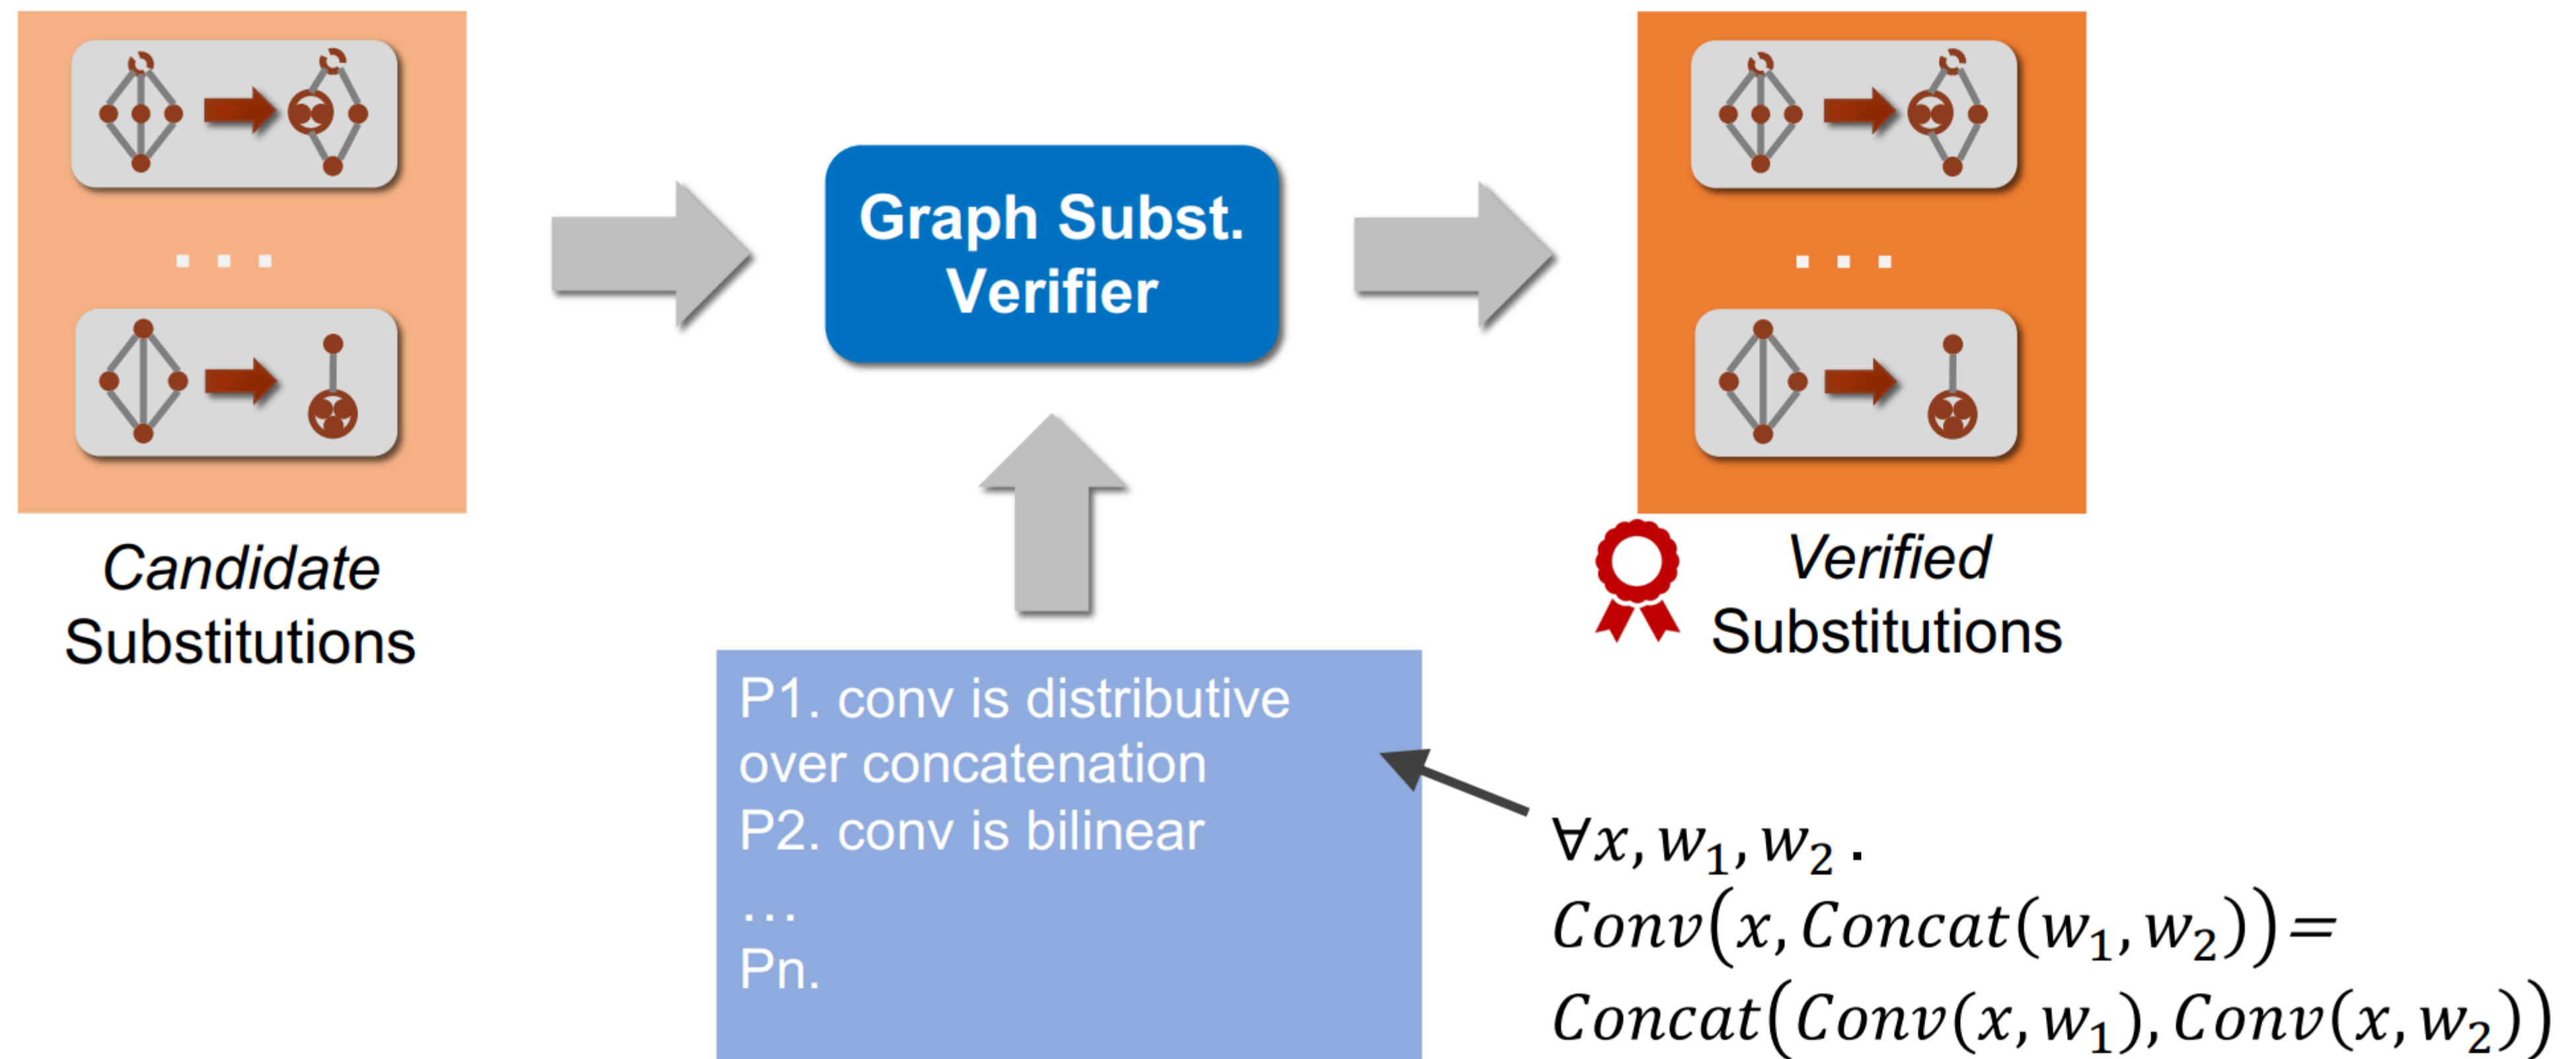
Common subgraph

734
substitutions

Can we trust graph substitutions?

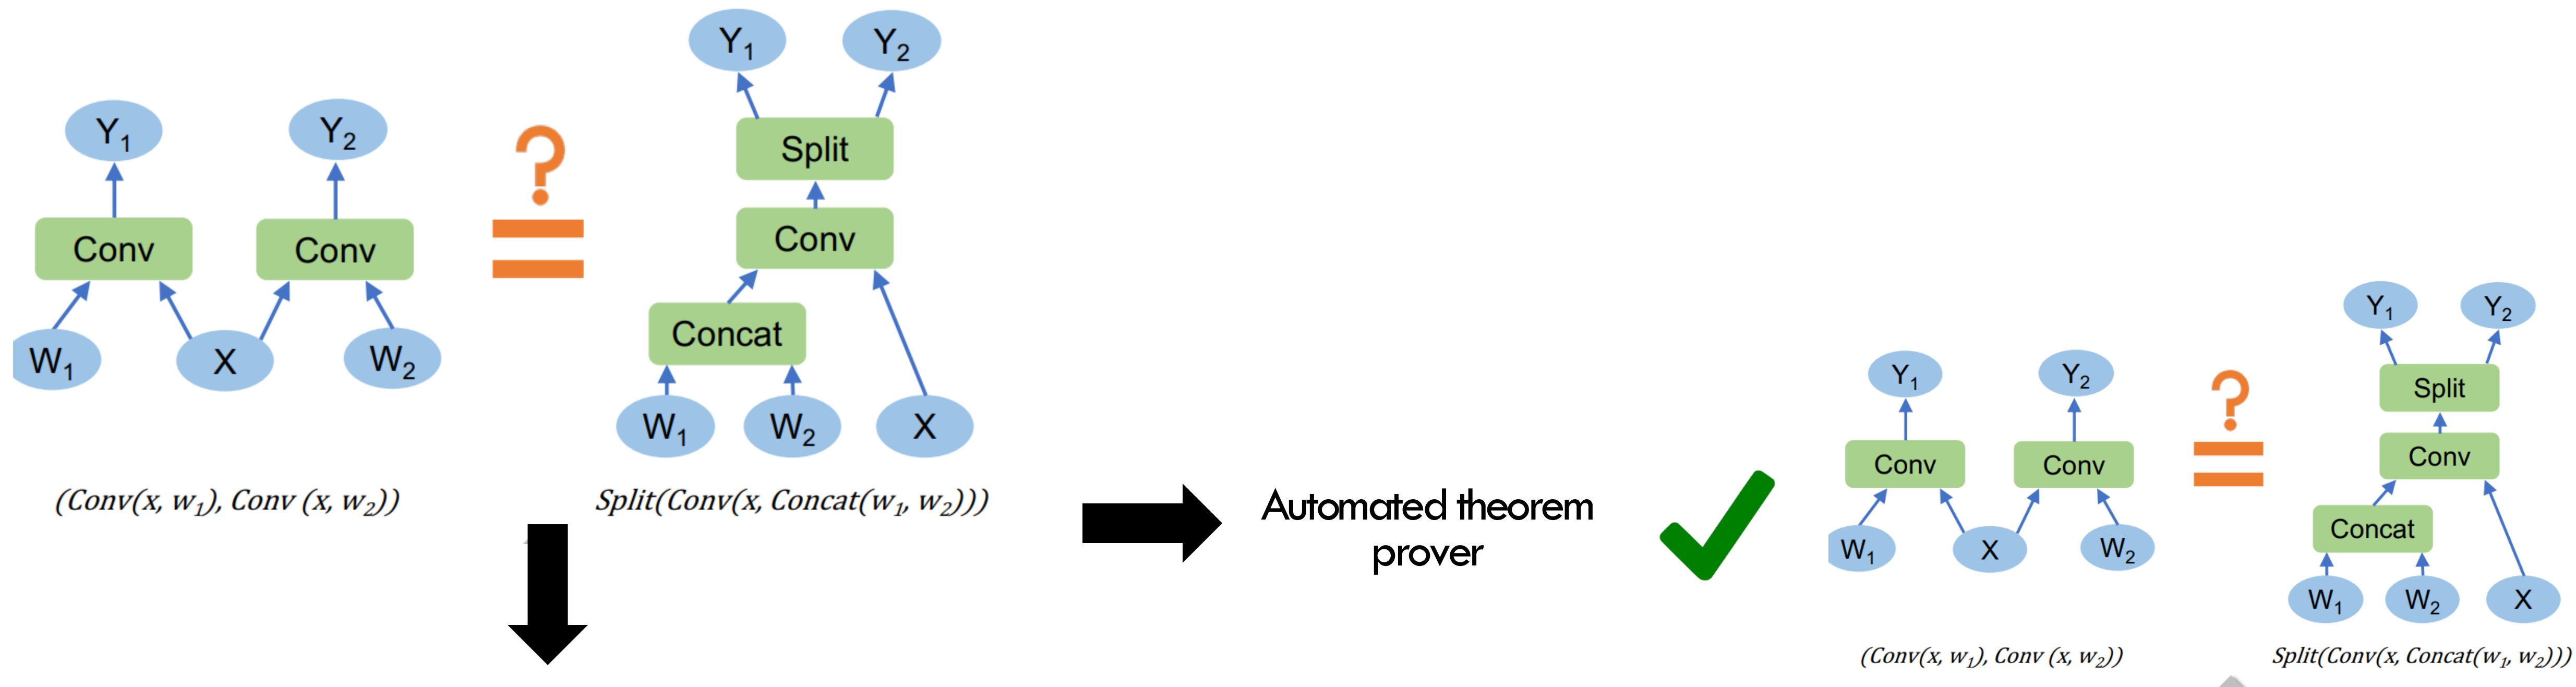
- We have $f(a) = g(b)$, $f(b) = g(b)$
 - But can we say: $f(x) = g(x)$ for $\forall x$
- We need to verify formally.

Substitution Verifier



Idea: writing specifications are easier than actually, conducting the optimizations

How to Verify



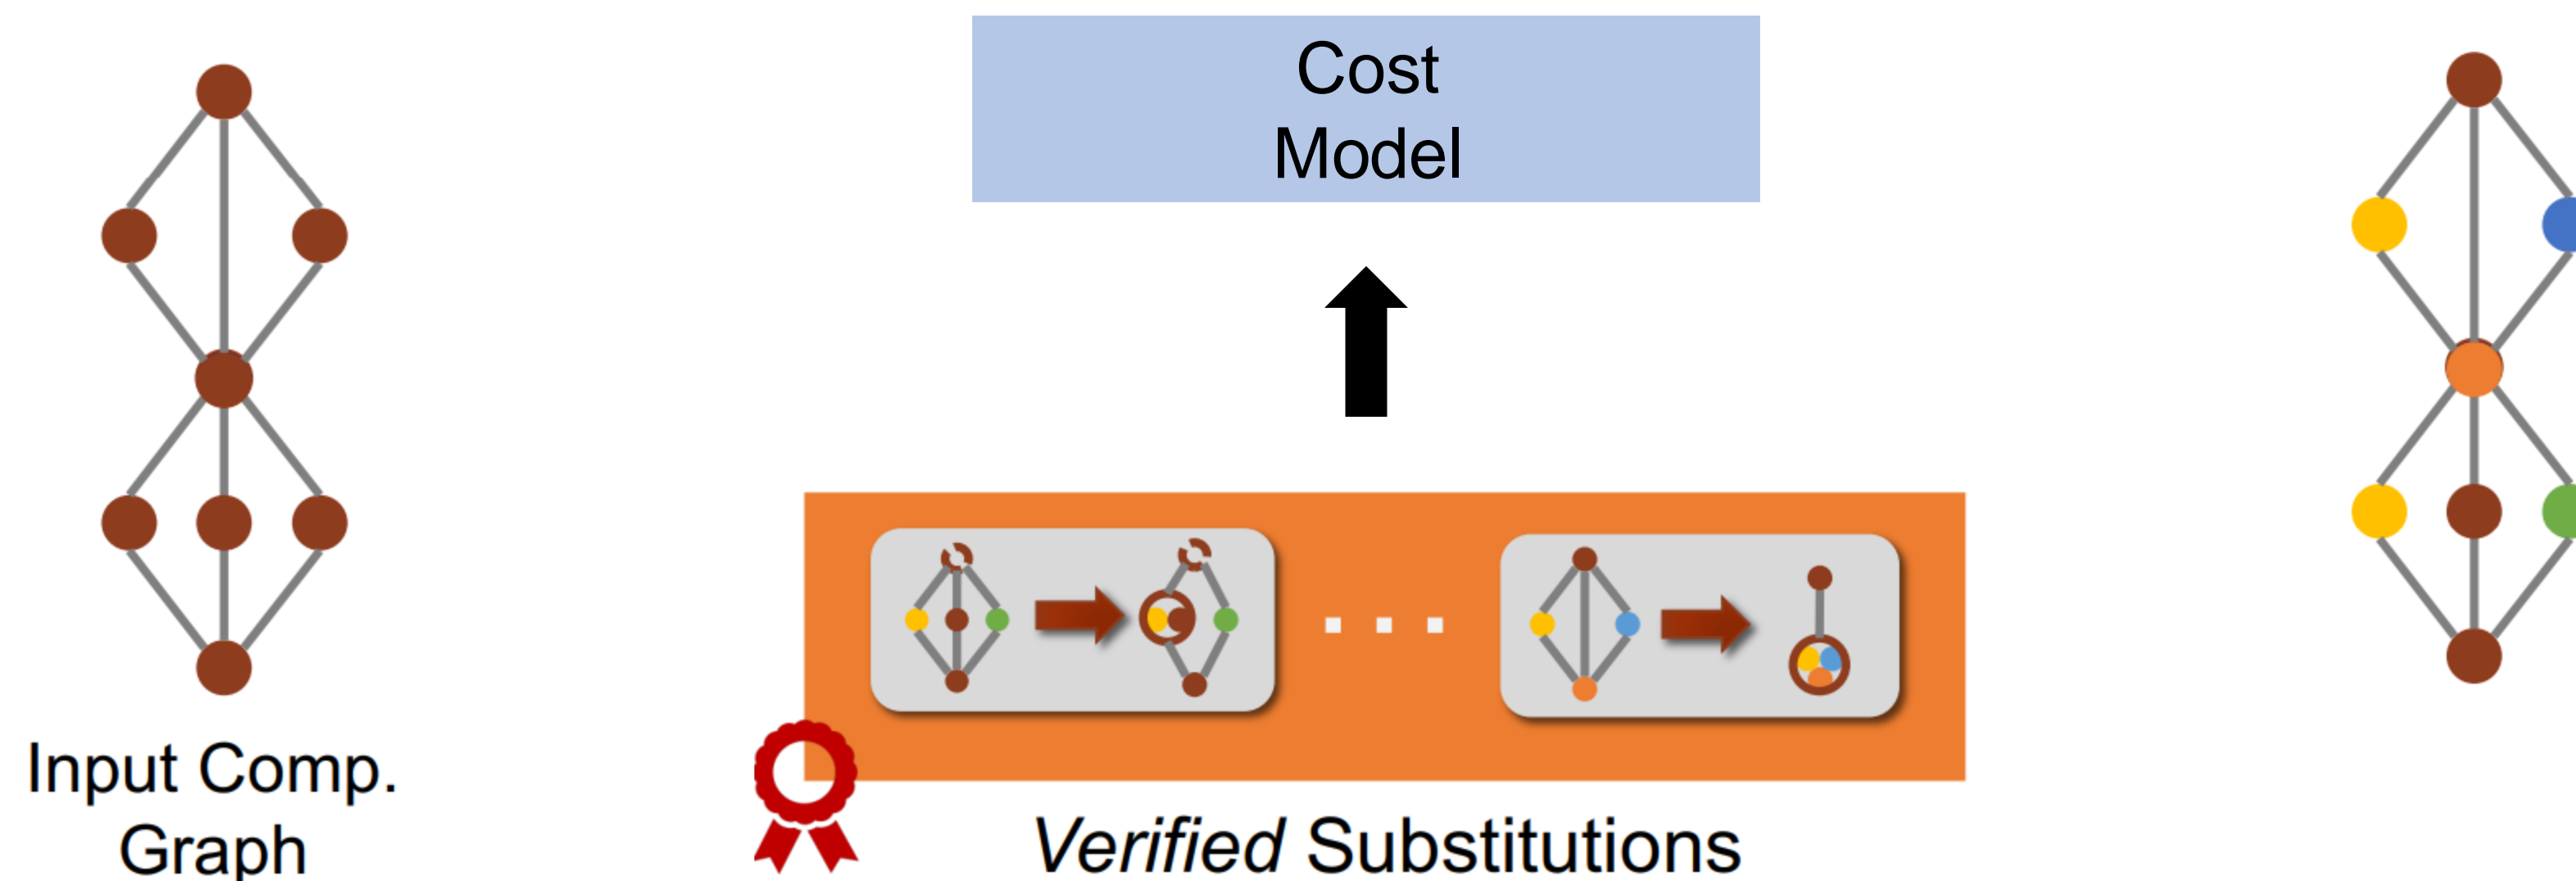
$\forall x, w_1, w_2 .$
 $(Conv(x, w_1), Conv(x, w_2))$
 $= Split(Conv(x, Concat(w_1, w_2)))$

P1. $\forall x, w_1, w_2 .$
 $Conv(x, Concat(w_1, w_2)) =$
 $Concat(Conv(x, w_1), Conv(x, w_2))$
 P2. ...


- Generating 743 substitutions = 5 mins
- Verify against 43 op specs = 10 mins
- Supporting a new op requires experts to write specs = 1400 LoC
 - vs. 53K LoC of manual optimization in TF

Incorporating substitutions

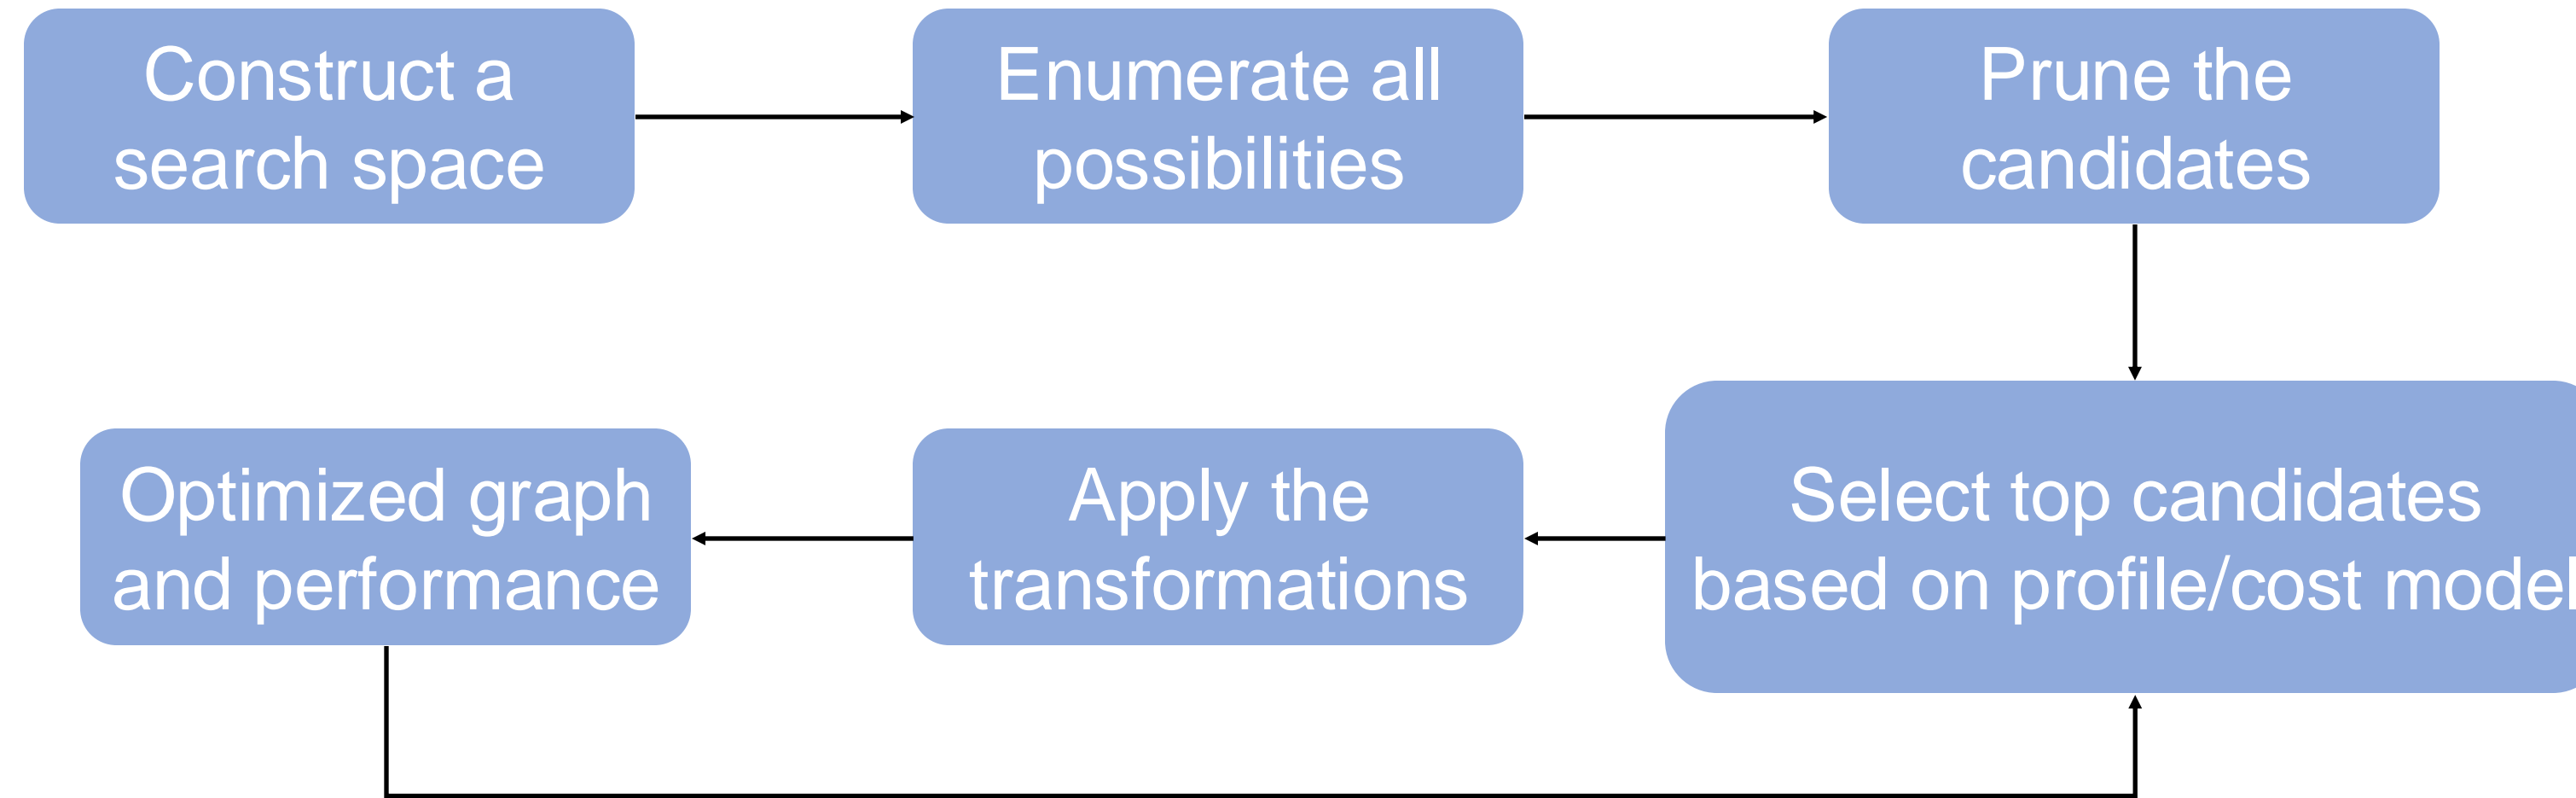
- Goal: apply verified substitutions to obtain an optimized graph
- Cost Model
 - Based on the sum of individual operator's cost
 - Profile each operator's cost on the target hardware
- Traverse the graph, apply substitutions, calculate cost, use backtracking



Performance (as of 2019)



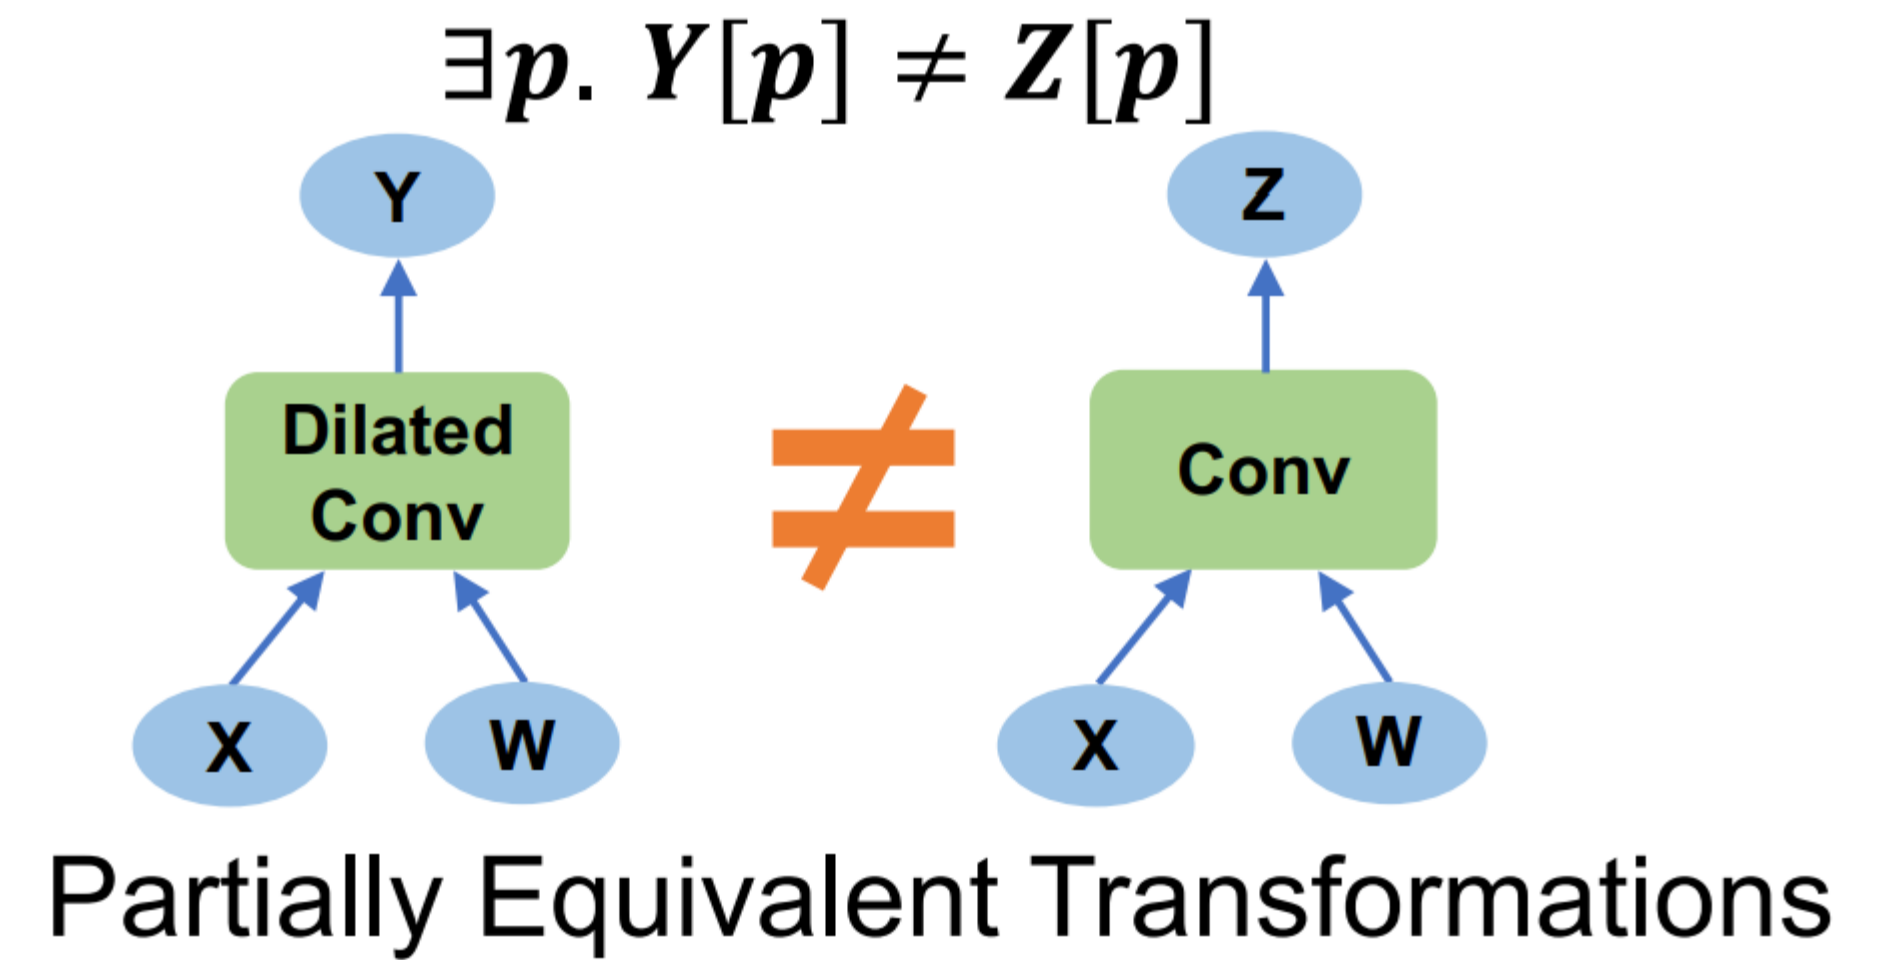

Summary of Graph Optimization



Limitations

- The best optimization is not covered by search space
- Search is too slow
- Evaluation of the resulting graph is too expensive
 - Limits your trial-and-error times

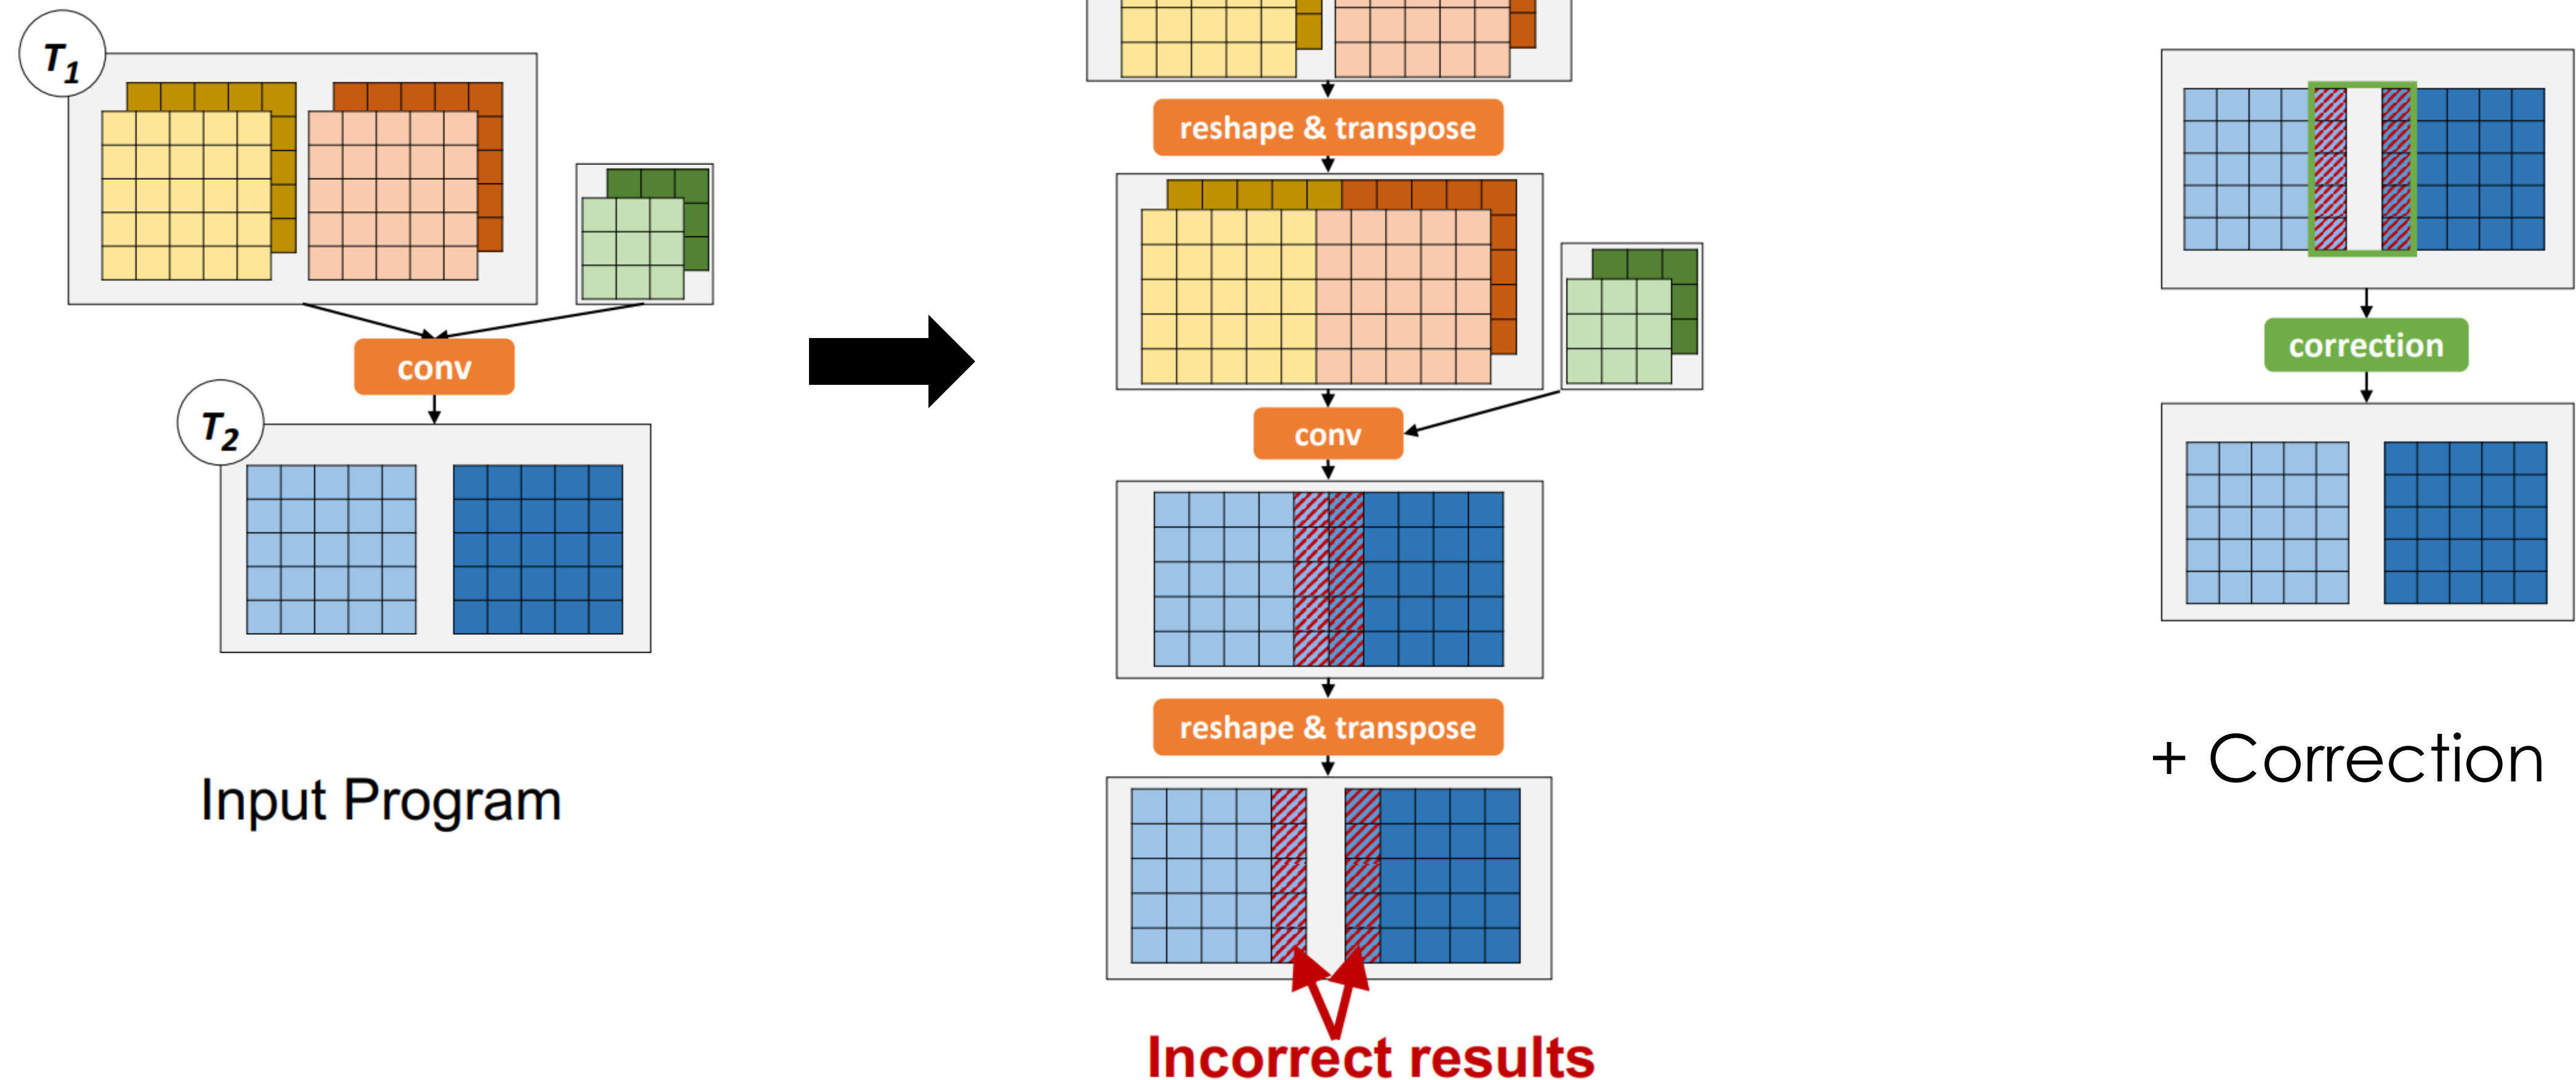
A Failure Example



- Math-equivalent
- Missing some optimization opportunities
- Better performance
- Not fully equivalent -> accuracy loss

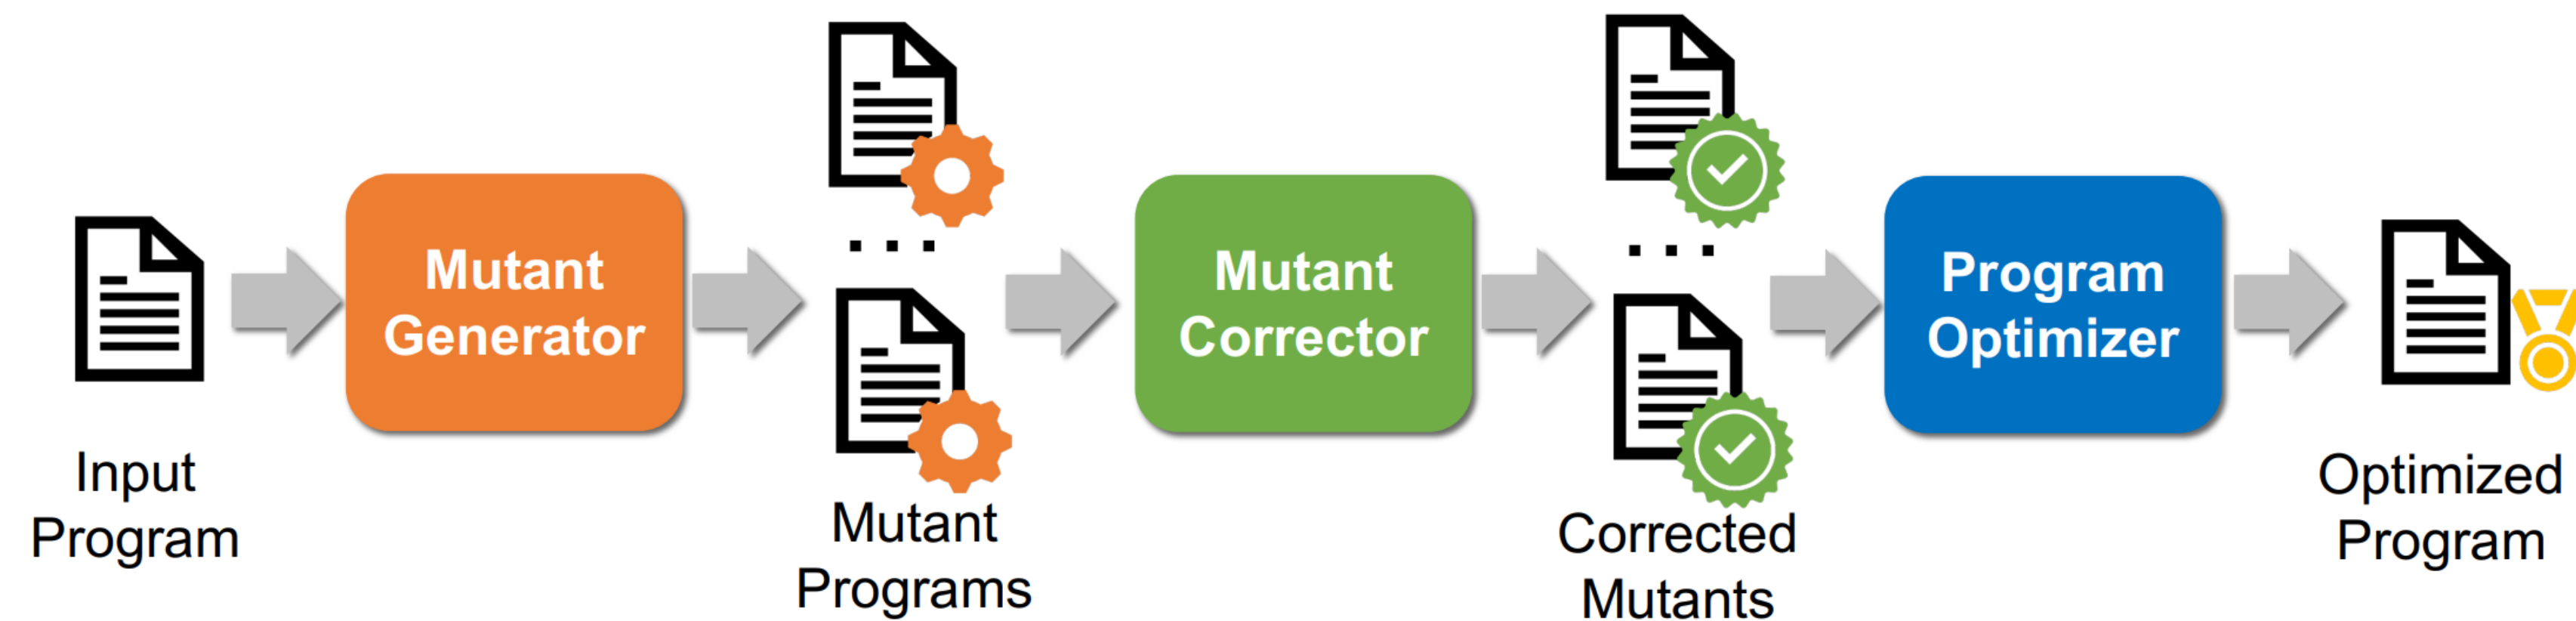
How about: exploit the larger space partially equivalent transformations for performance while still preserve correctness?

Motivating Example



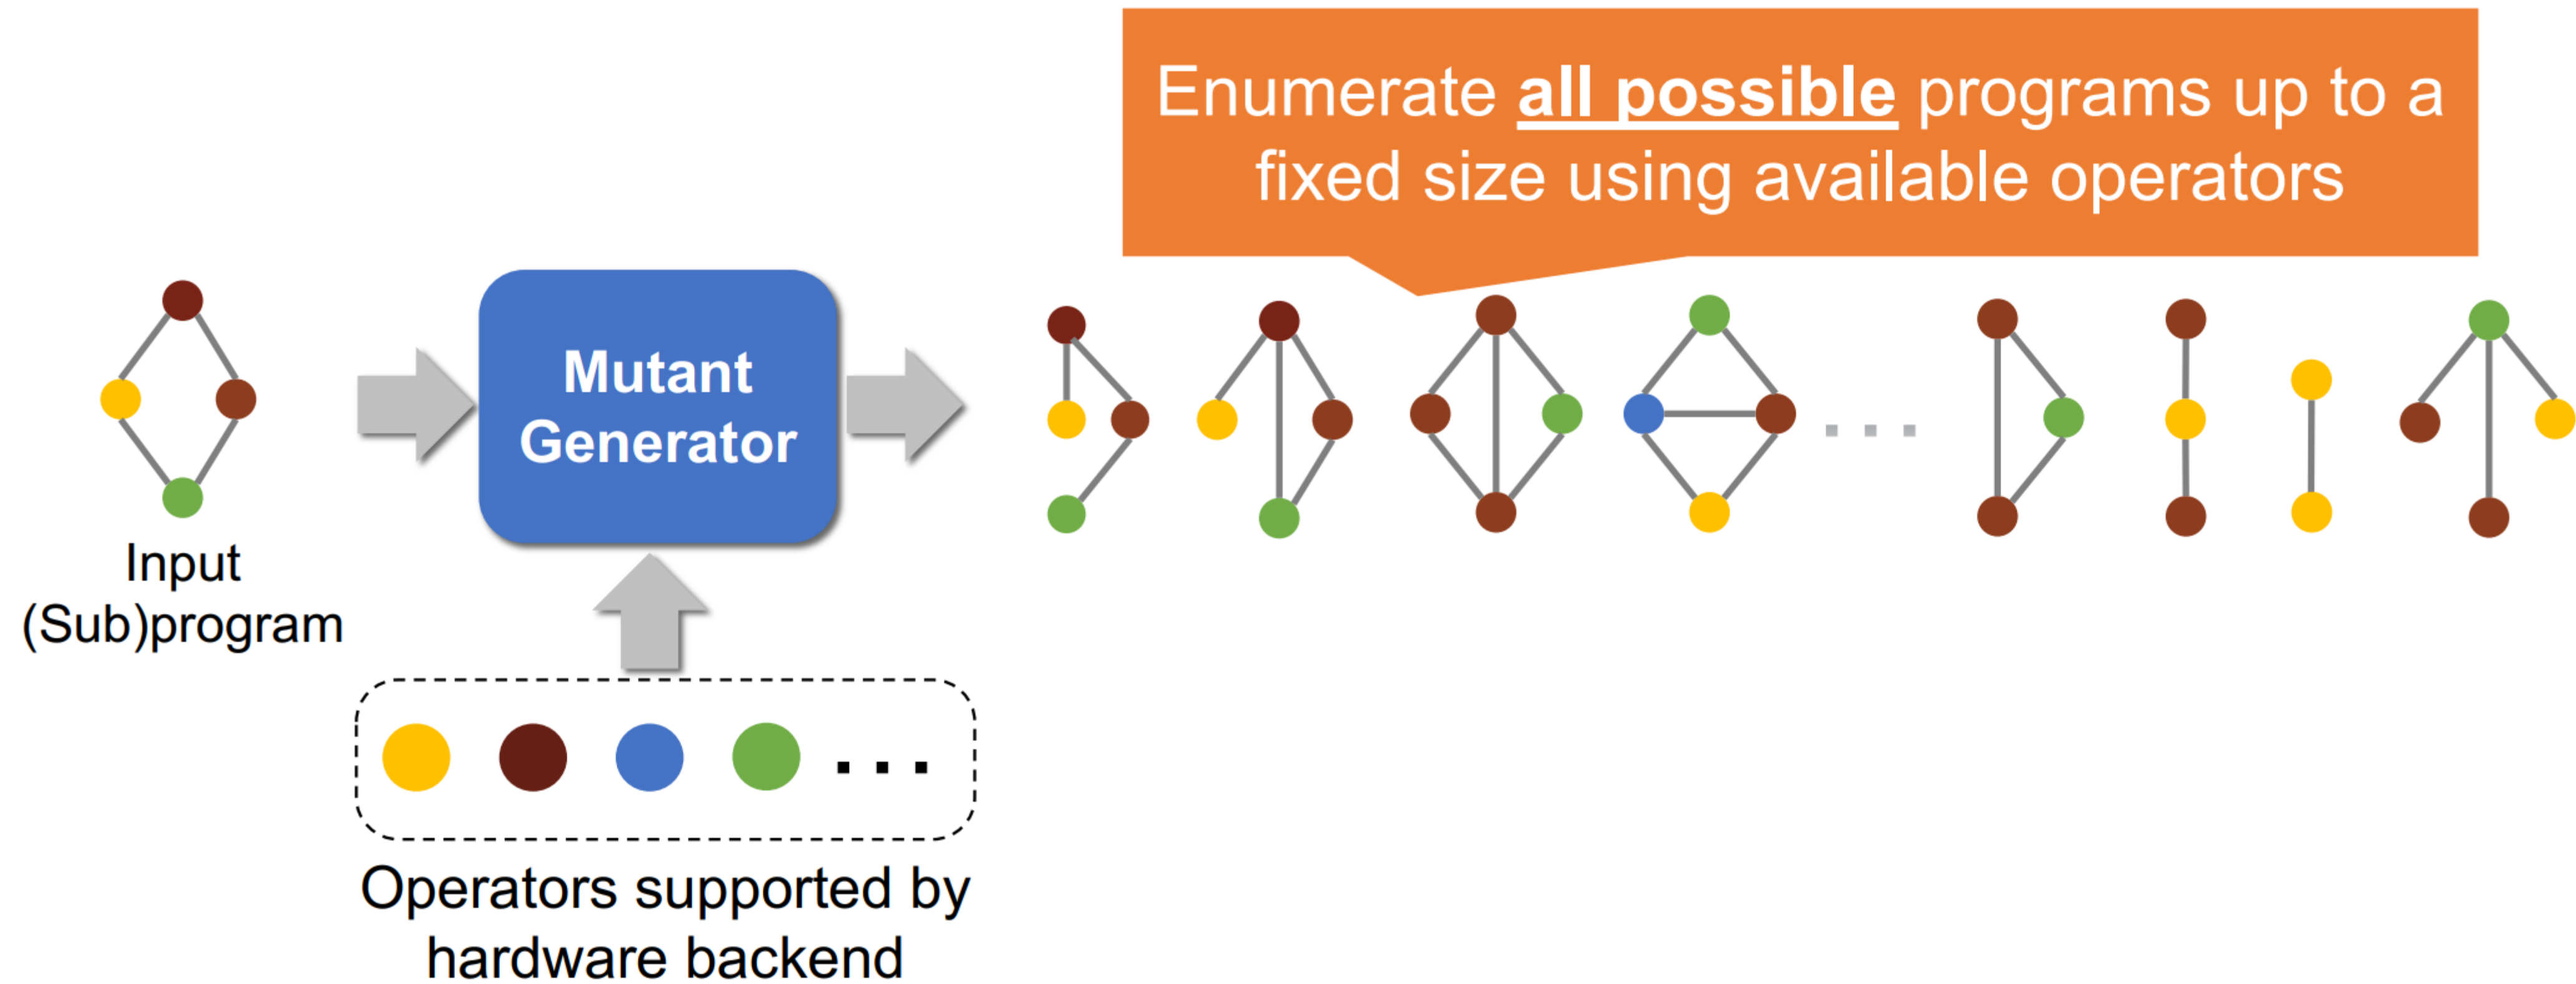
- Partial equivalent transformations + correction yield 1.2x speedup
- Which would otherwise be impossible in fully equivalent transformations space

Partially Equivalent Transformations

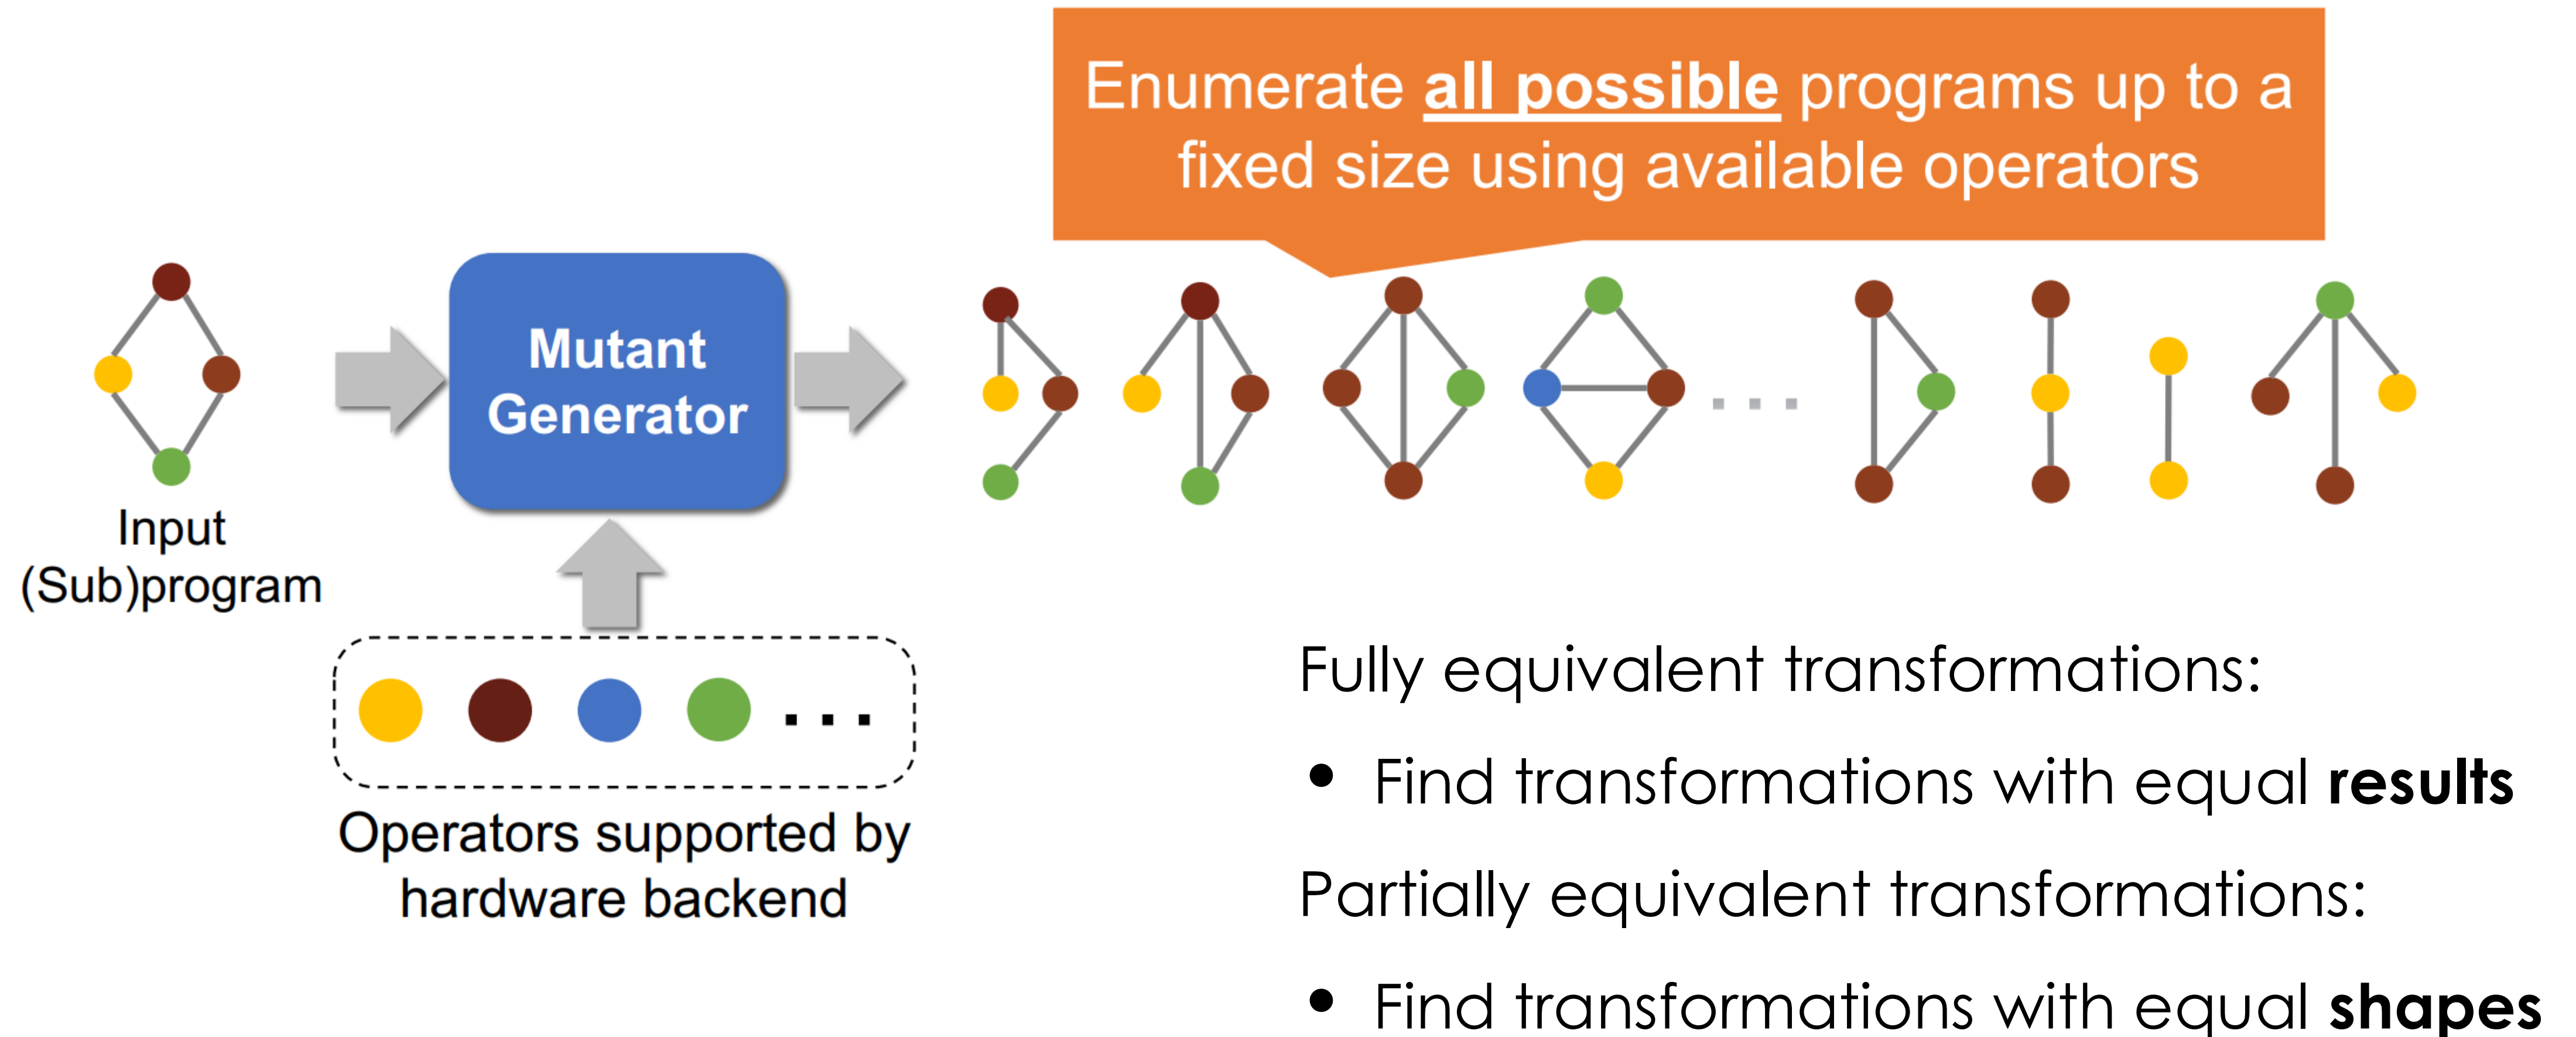


- How to mutate?
- How to correct?

Mutant Generator: Step 1

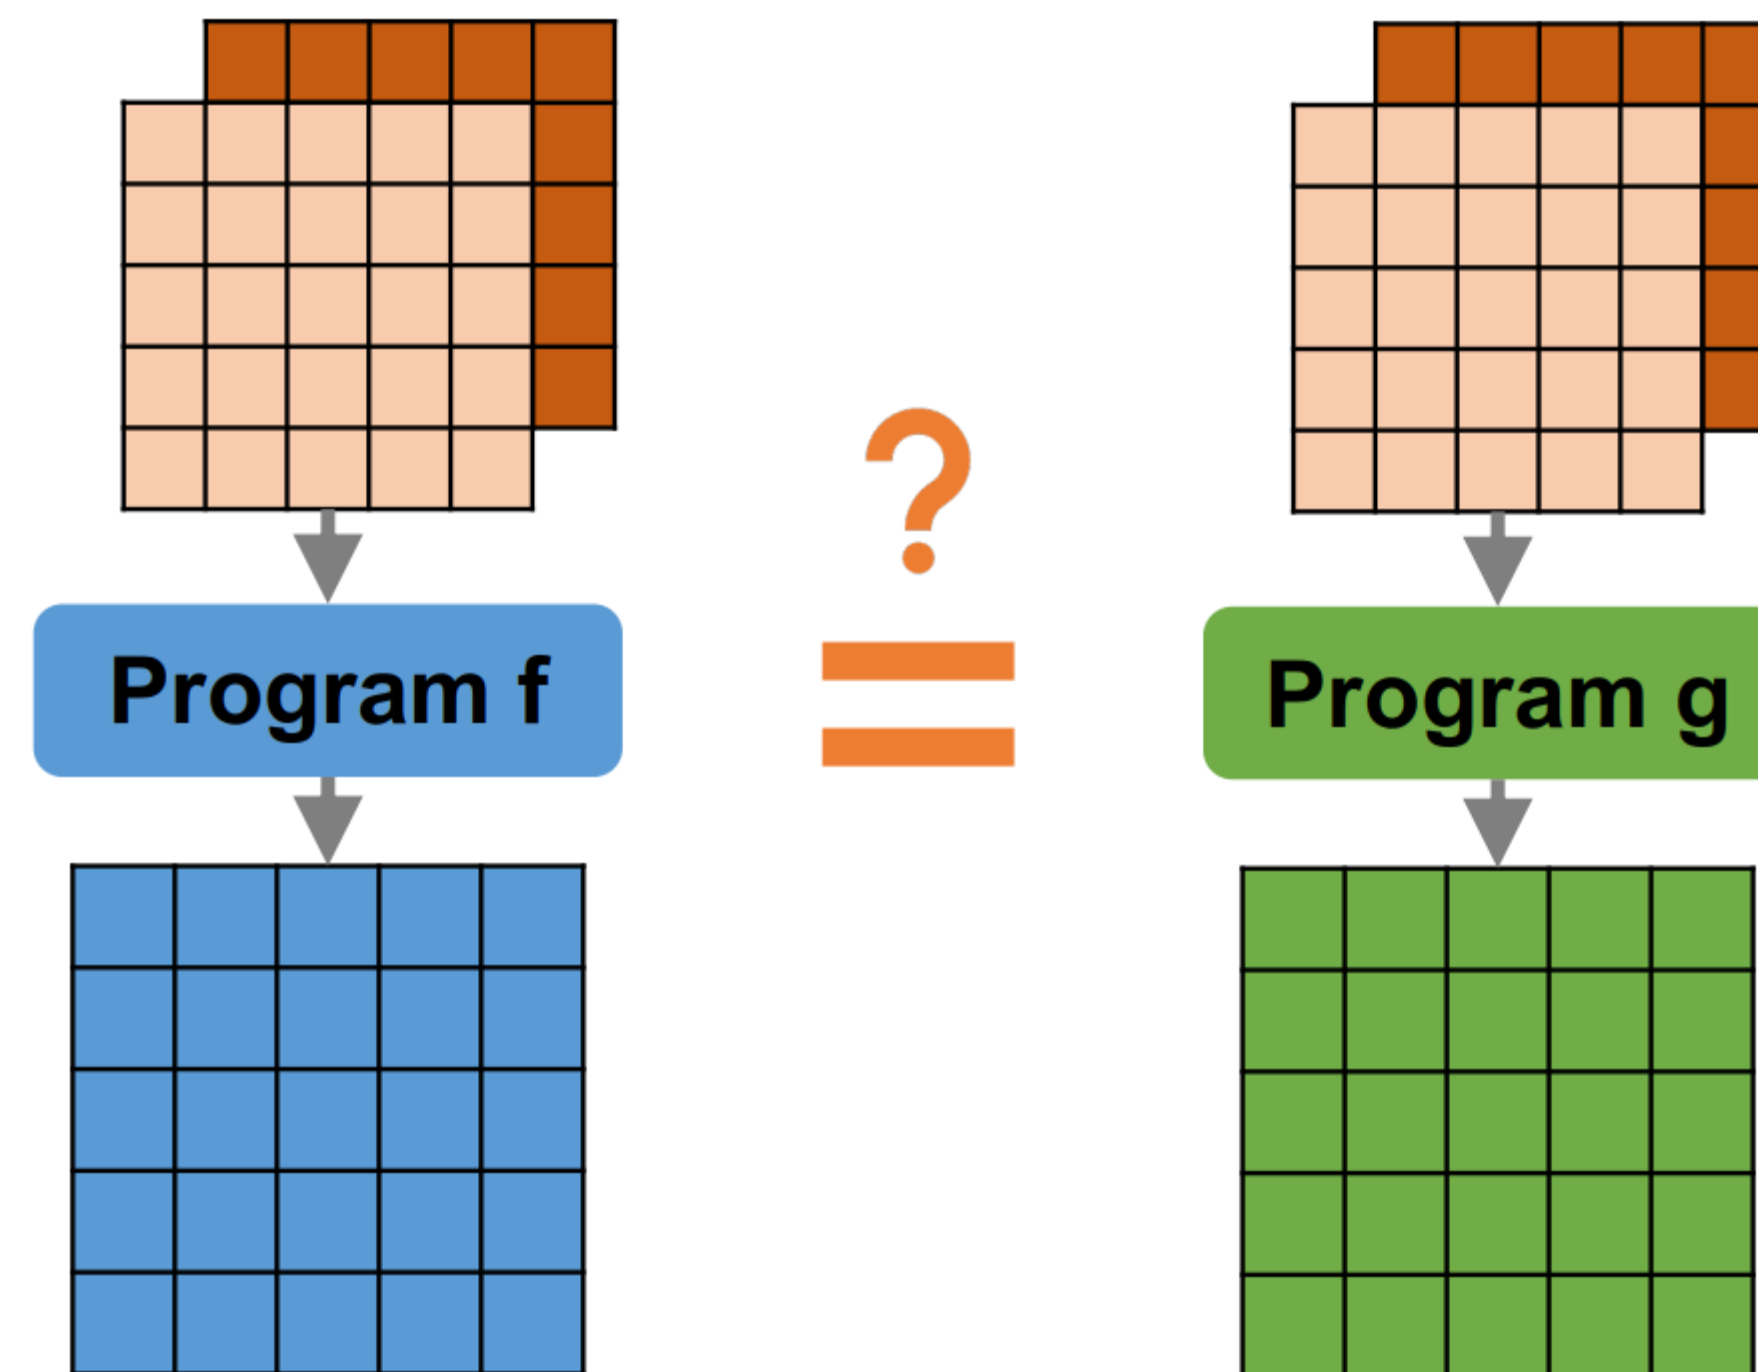


Mutant Generator: Step 2



How to Detect and Correct?

- Which part of the computation is not equivalent?
- How to correct the results?



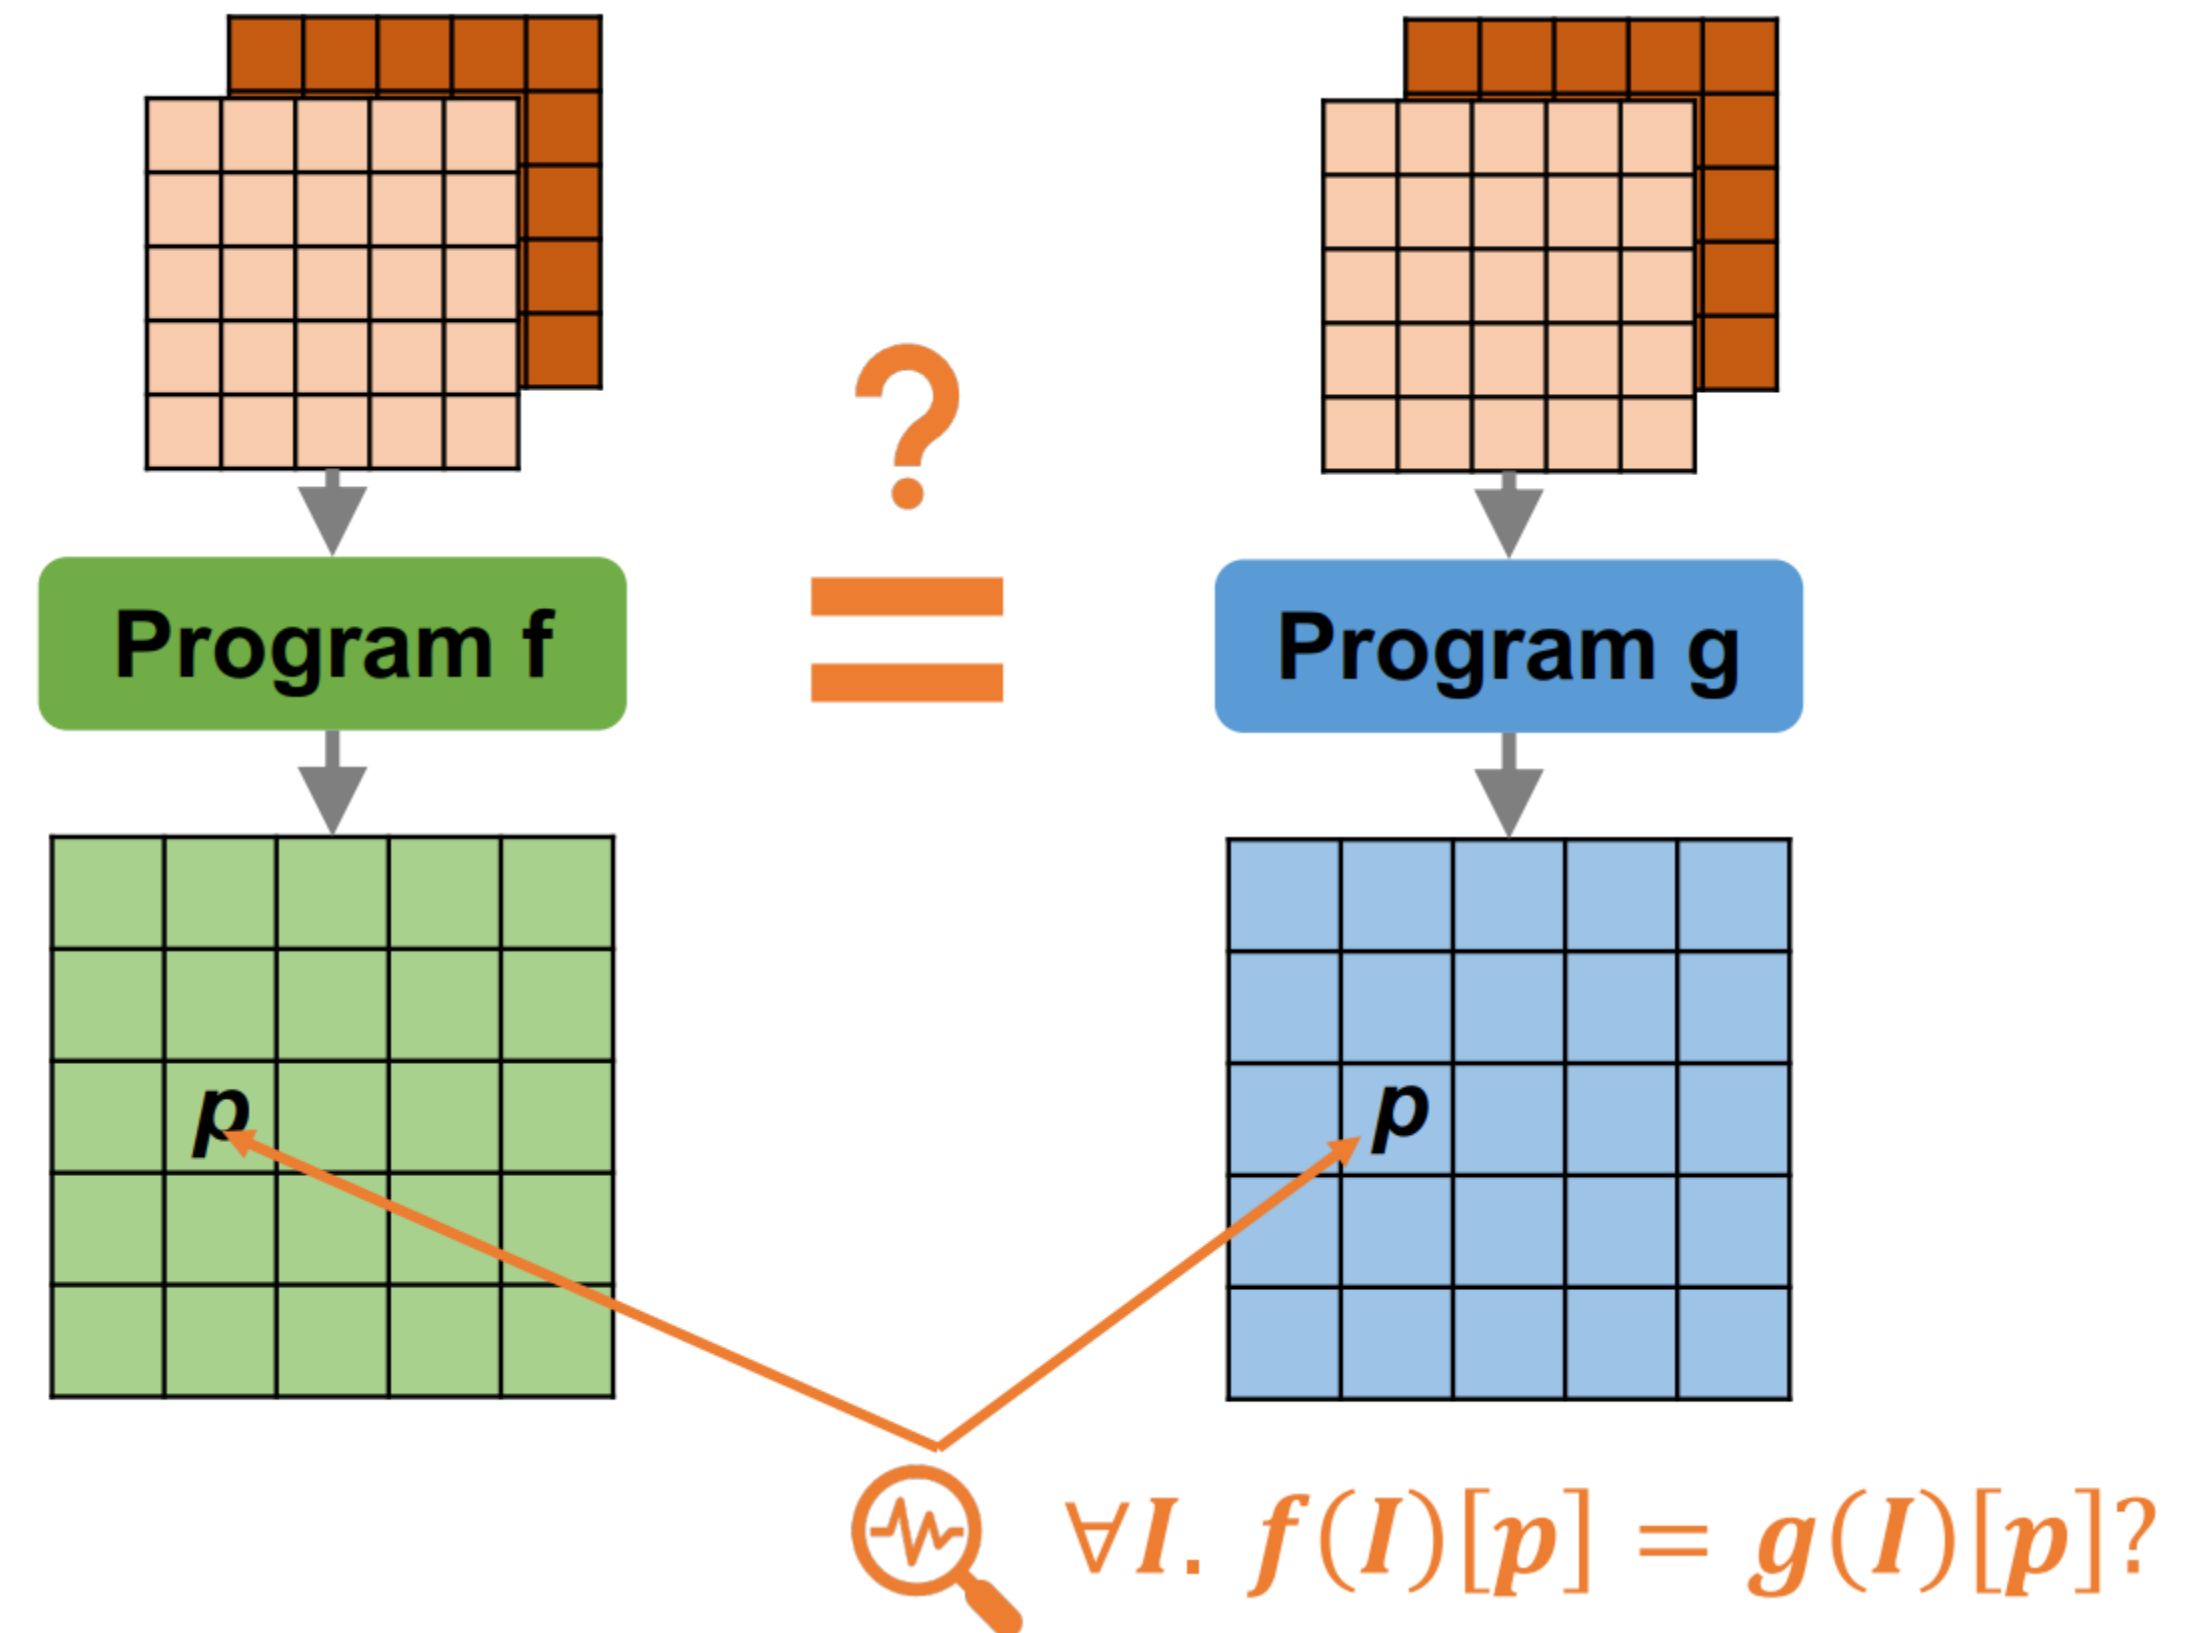
By Enumeration

- For each possible input I
 - For each position p
 - Check if $f(I)[p] == g(I)[p]$

- Complexity $O(m \times n)$:

- m : possible inputs
- n : output shape

- How to reduce enumeration effort?
 - Reduce m and n




How to reduce n?

- Can we just check out a few (or even just one) position at $f(I)[p]$ and assert the (in-)correctness?
- Answer: Yes for 80% of the computation
- Reason: Neural nets computation are mostly Multi-Linear
- Define Multi-linear: f is multi-linear if the output is linear to all inputs
 - $f(I_1, \dots, X, \dots, I_n) + f(I_1, \dots, Y, \dots, I_n) = f(I_1, \dots, X + Y, \dots, I_n)$
 - $\alpha f(I_1, \dots, X, \dots, I_n) = f(I_1, \dots, \alpha X, \dots, I_n)$

How to reduce n

- Theorem 1: For two Multi-linear functions f and g , if $f=g$ for $O(1)$ positions in a region, then $f=g$ for all positions in the region
- Implications: only need to examine $O(1)$ positions for each region
- Reduce $O(mn) \rightarrow O(m)$

Group all output positions with an identical summation interval into a region

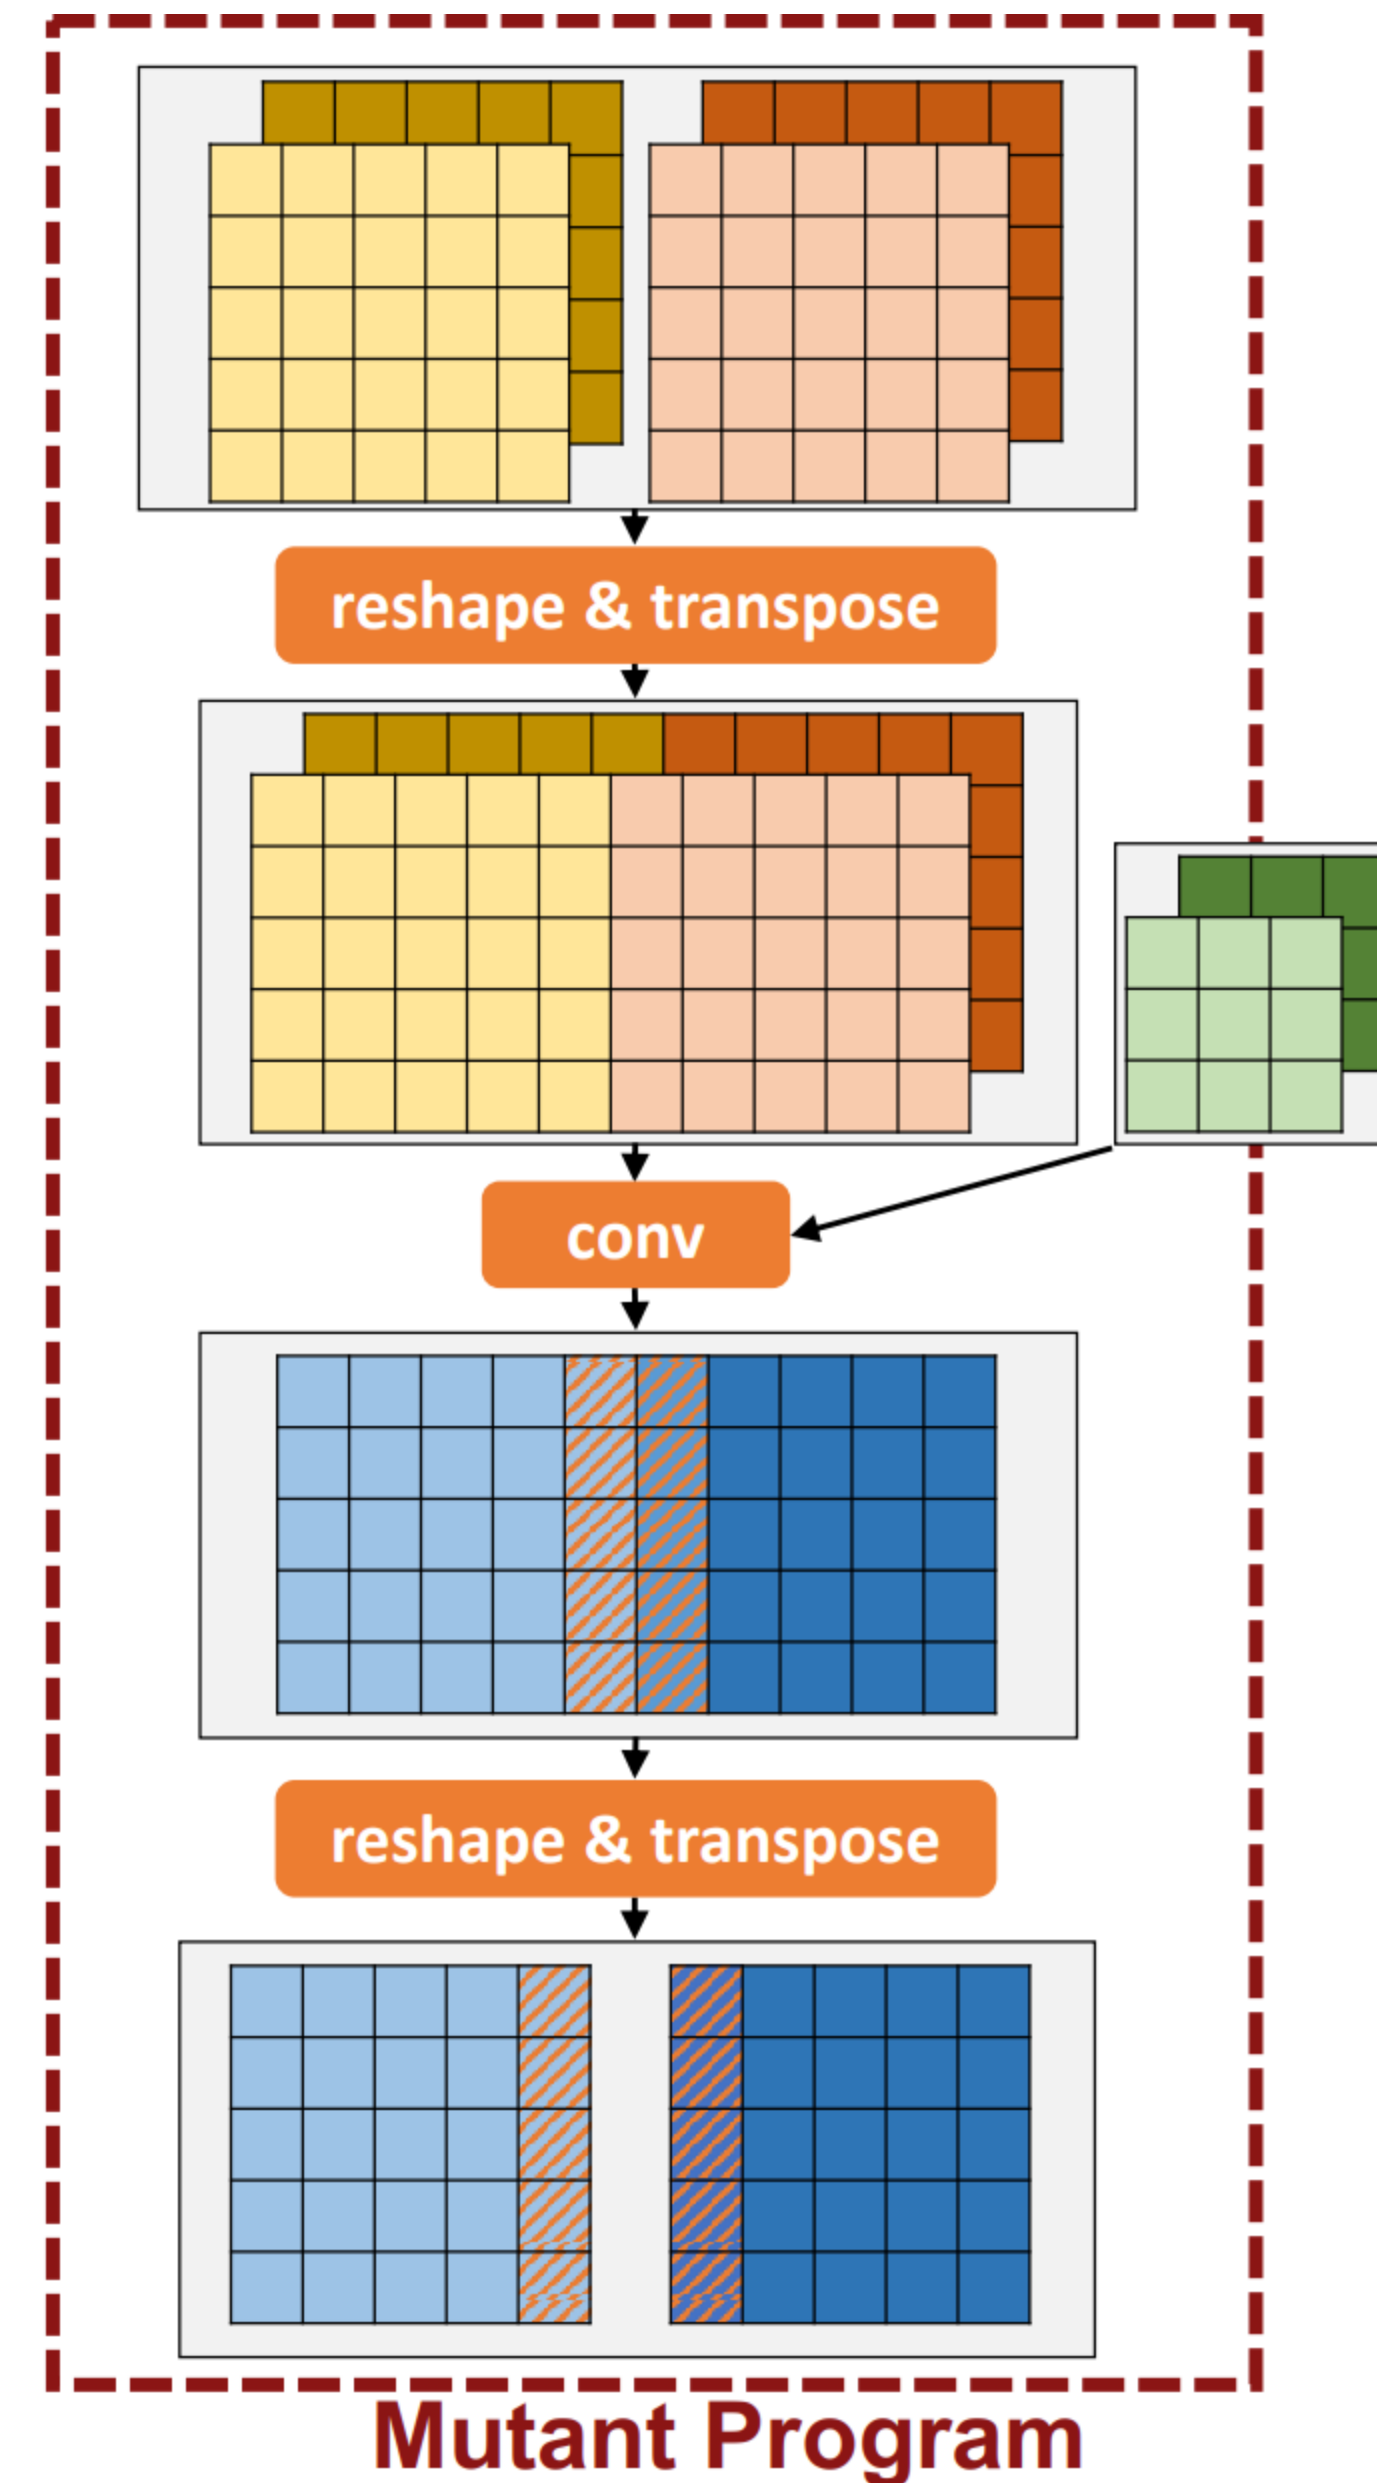


How to reduce m?

- Theorem 2: if $\exists I, f(I)[p] \neq g(I)[p]$, then the probability that f and g give identical results on t random inputs is $\left(\frac{1}{2^{31}}\right)^t$
- Implications: Run t random tests with random input, and if all t passed, it is very unlikely f and g are inequivalent
- $O(mn) \rightarrow O(m) \rightarrow O(t)$ ($t \ll m$)


Correct the Mutant

- Goal: quickly and efficiently correcting the outputs of a mutant program




Correct the Mutant

- Goal: quickly and efficiently correcting the outputs of a mutant program
- Step 1: recompute the incorrect outputs using the original program

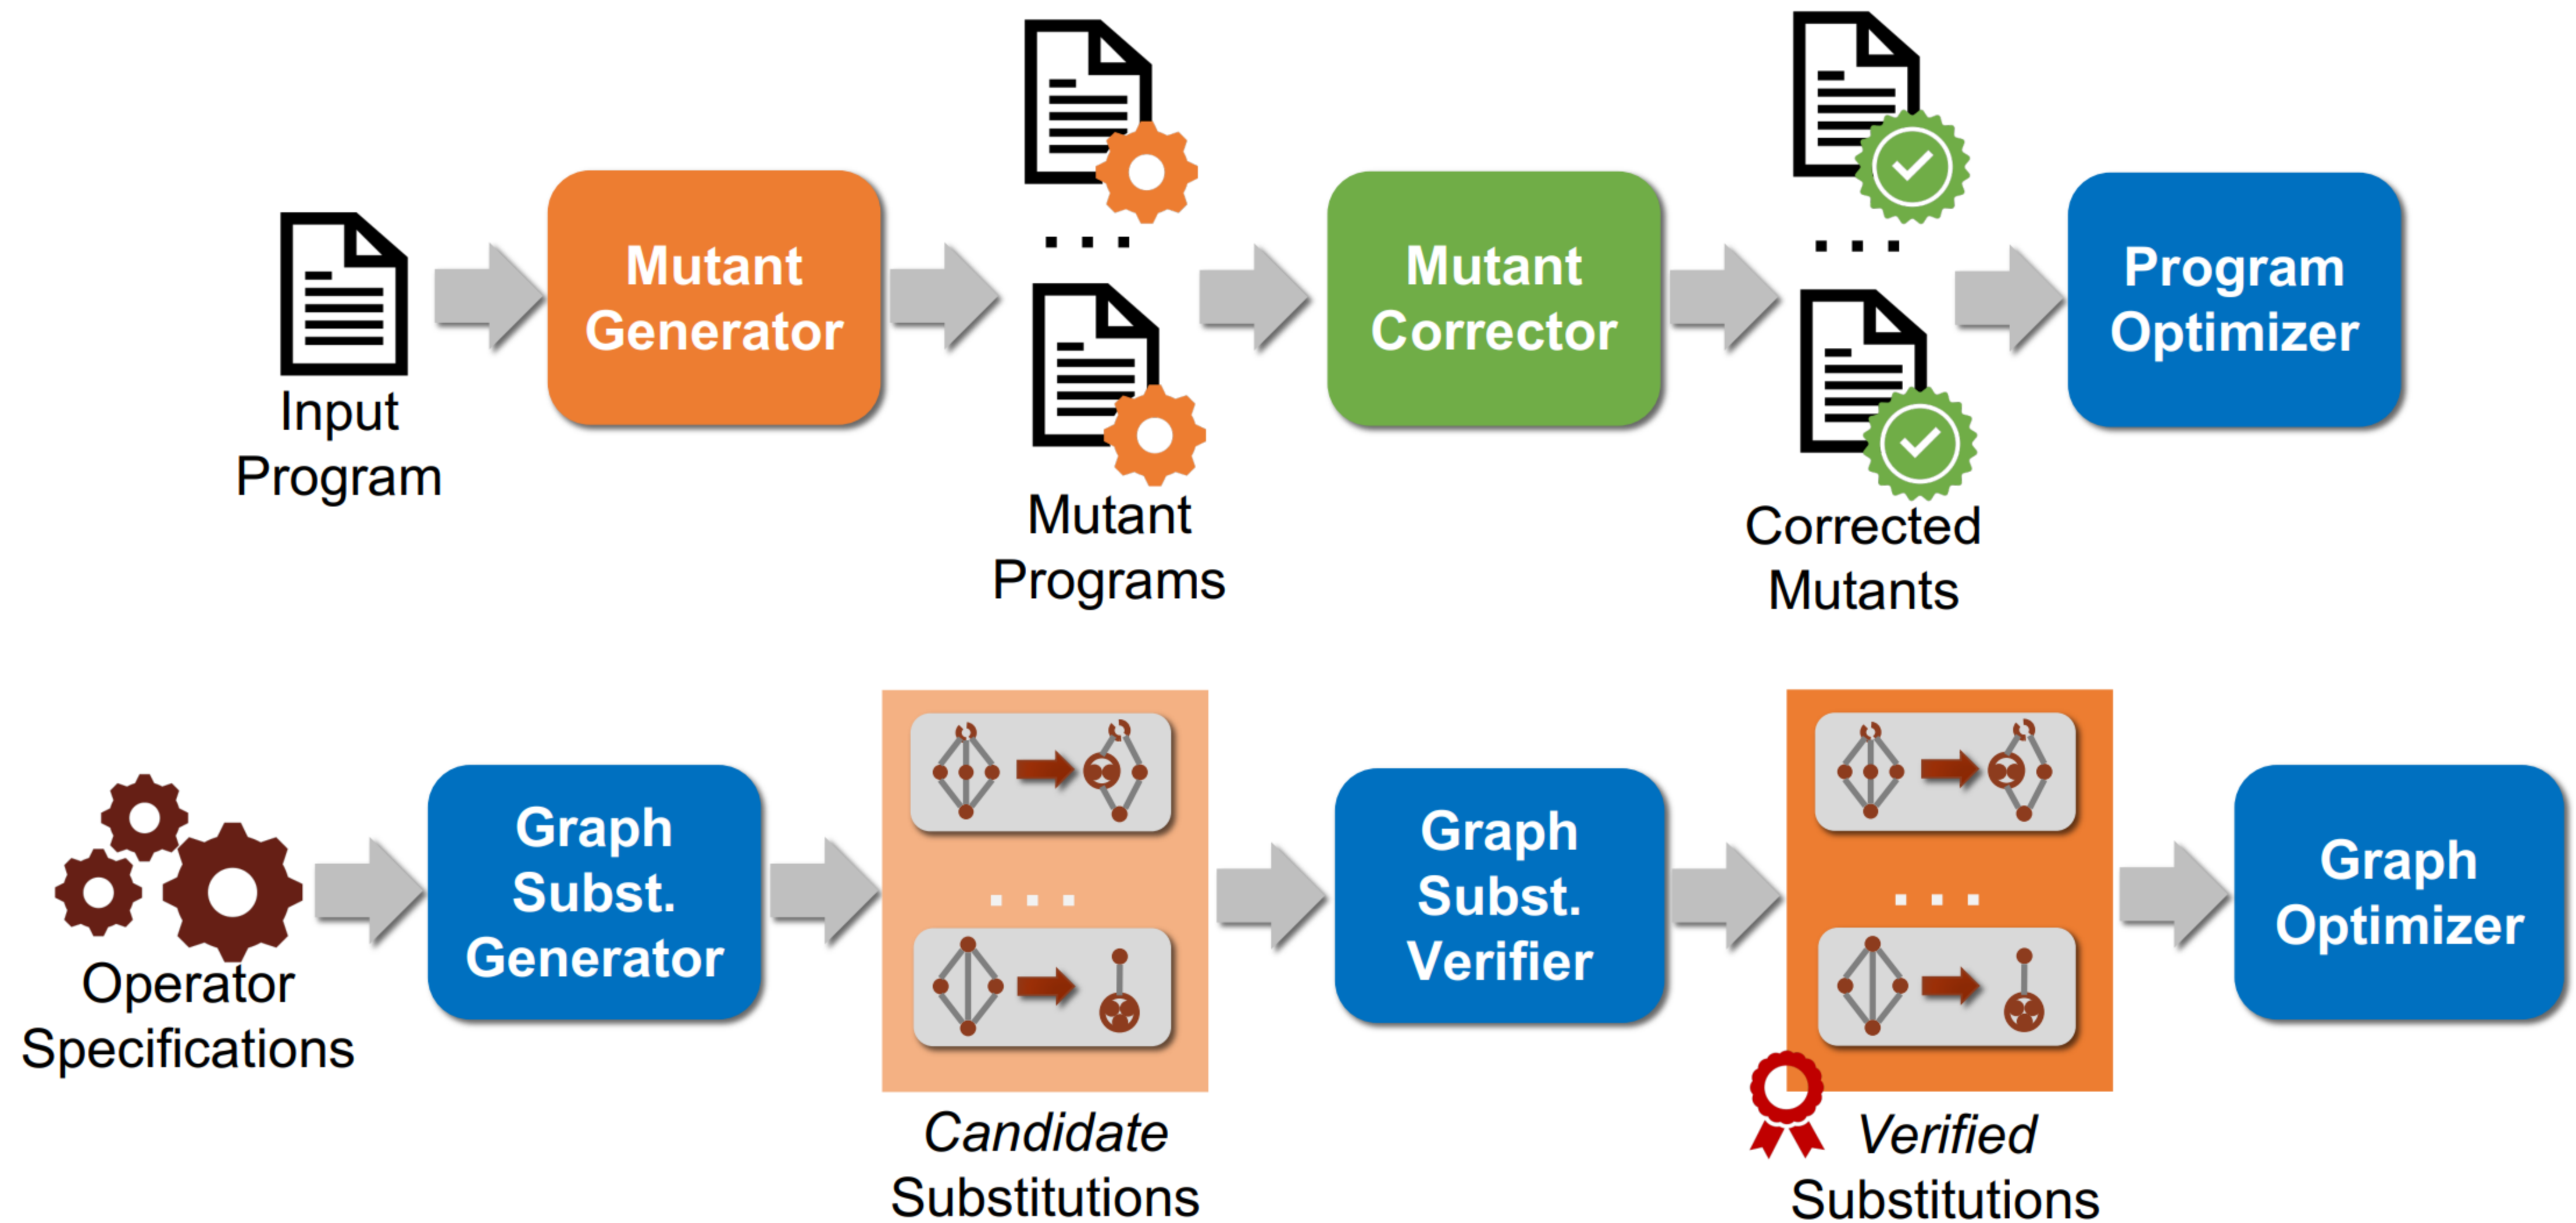


Correct the Mutant

- Goal: quickly and efficiently correcting the outputs of a mutant program
- Step 1: recompute the incorrect outputs using the original program
- Step 2: opportunistically fuse correction kernels with other operators



Recap




Summary & Questions to discuss

- Fully equivalent transformations vs. Partial
 - How to define search space
 - How to prune search space
 - How to verify & correct
 - How to apply to the ML graph optimization

ML Compiler Retrospective

Q: why the community shifts away from compiler

500+ compiler papers are written during


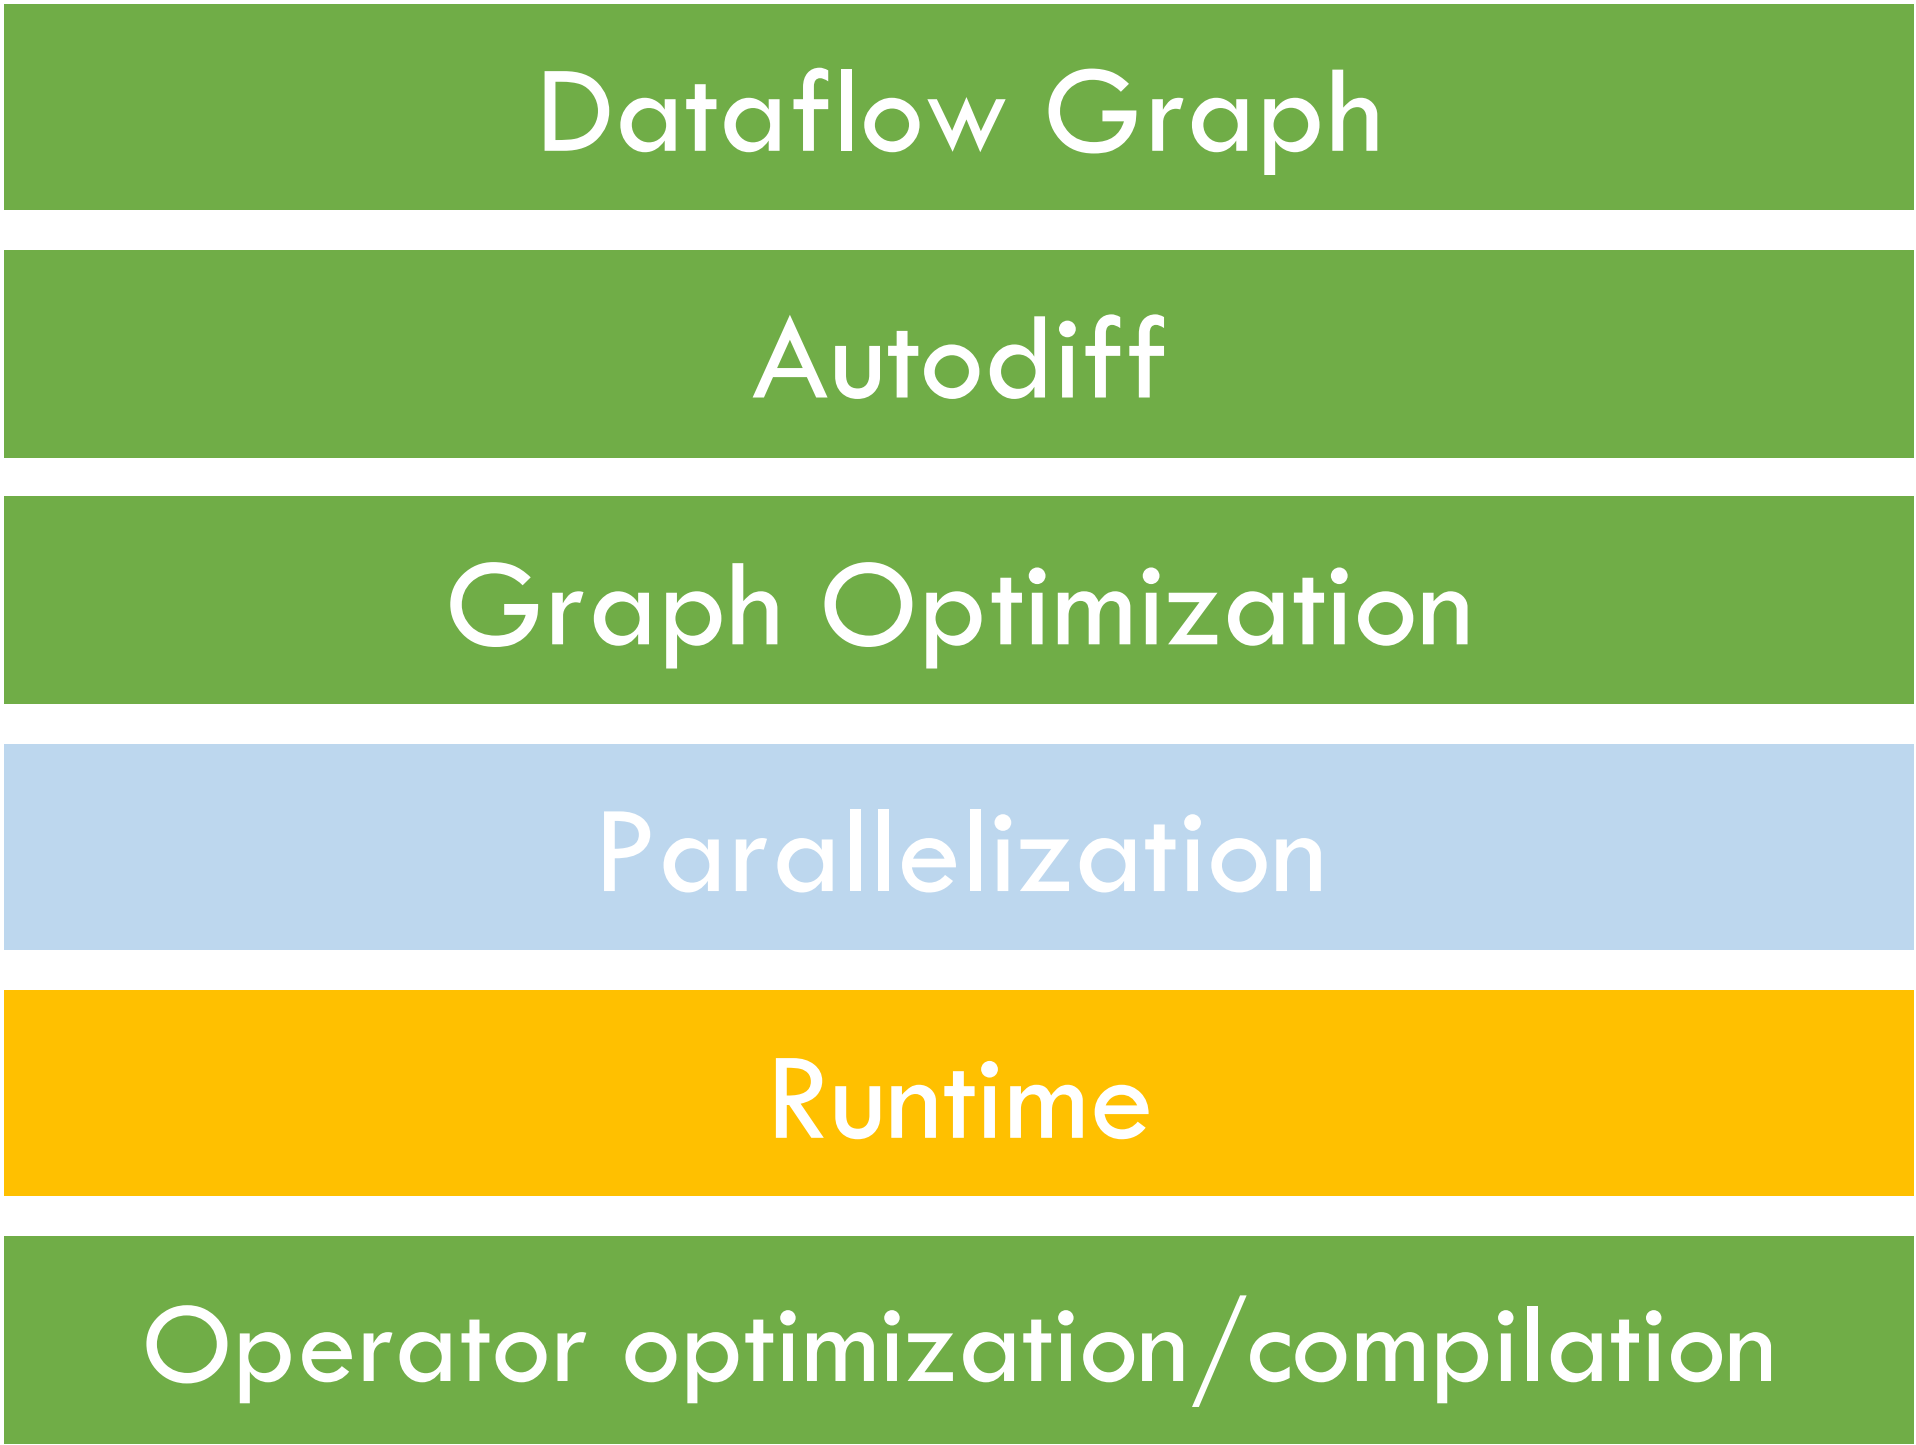
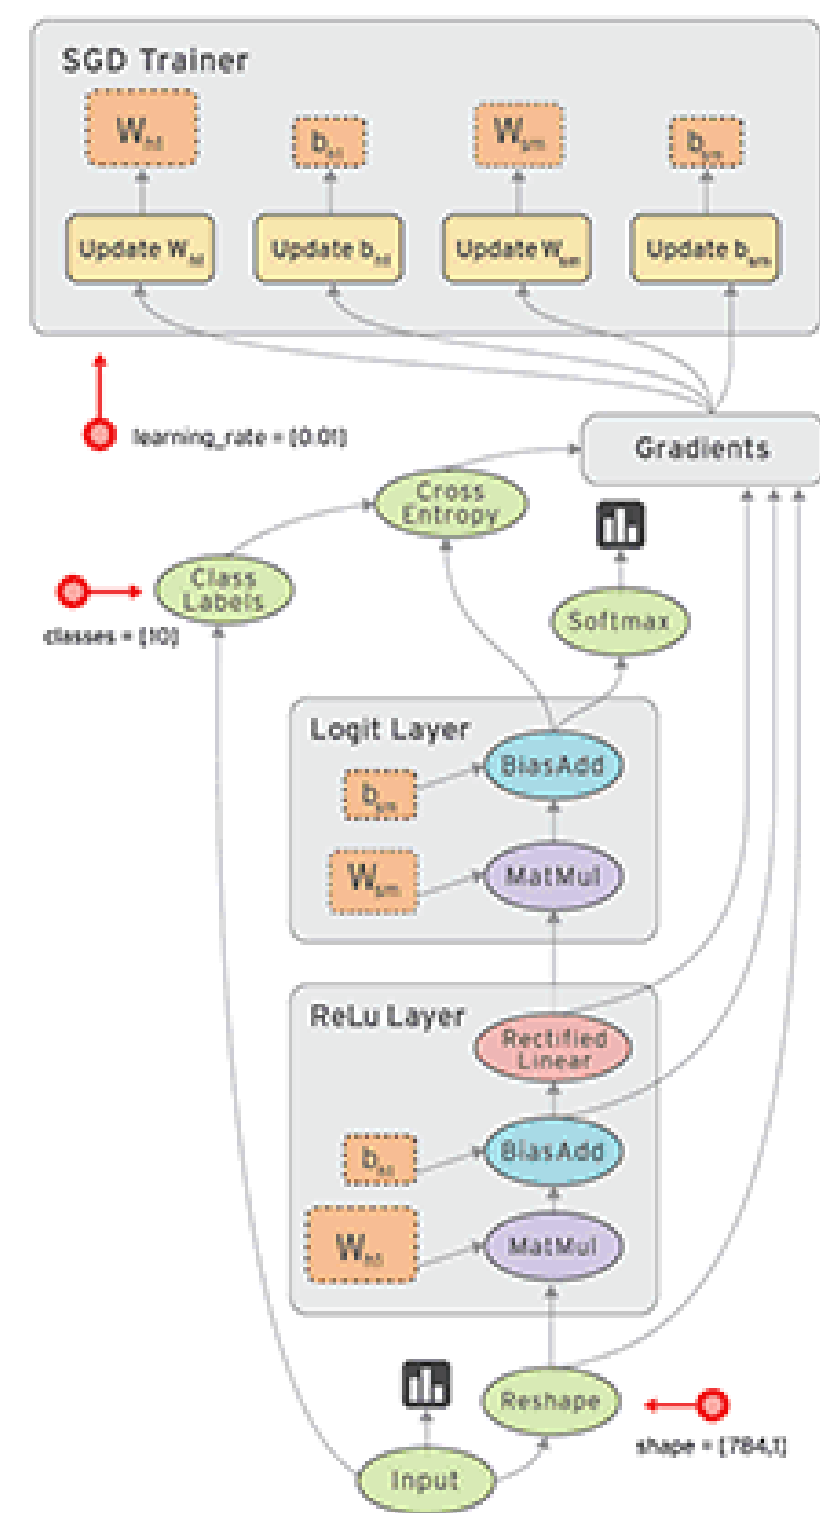


More Compiler / Graph Optimization



- Guest Speaker: Tianqi Chen
- A.k.a.: GOAT of MLSys
- Inventor of: XGBoost, TVM, MLC-LLM
- Date: Feb. 6
- Topic: Machine Learning
Compilation

Big Picture: where are we



Next: Runtime

- “Batching”
- Checkpointing and rematerialization
- Swapping
- Quantization, Mixed precision, and Pruning